

# [Chap.2-3] Representing and Manipulating Information

Young Ik Eom ([yieom@skku.edu](mailto:yieom@skku.edu), 031-290-7120)

Distributing Computing Laboratory

Sungkyunkwan University

<http://dclab.skku.ac.kr>



# Contents

- Introduction
- Information storage
- Integer representations
- Integer arithmetic
- Floating point
- Summary

## ■ The problem

- How to represent fractional values with finite # of bits?
  - 0.1
  - 2.718281828...
  - 3.14159265358979323846264338327950288...
  
- How to represent very small (close to 0) numbers ( $|V| \ll 1$ ) and very large numbers ( $|V| \gg 0$ )
  - 0.000...01
  - $0.123 \times 10^{-60}$
  - $1.234 \times 10^{90}$

# Floating Point

## ■ The problem

- 3 fields for representation of FP numbers
  - Sign
  - Exponent
  - Fraction (significand, mantissa)



# Floating Point

## ■ Warming up

### ■ Example-1)

- $0.1_{(10)} = 0.0001100110011[0011]_{(2)}$
- $0.110011001100... \times 2^{-3} \quad (\Leftrightarrow 0.00011001100... \times 2^0)$
- (s, exp, frac) with 16-bits FPN = 0 11101 1100110011

Excess-16 exp:  $11101 + 10000 = 01101$

### ■ Example-2)

- $5.1_{(10)} = 101.0001100110011[0011]_{(2)}$
- $0.10100011001100110011... \times 2^3$
- (s, exp, frac) with 16-bits FPN = 0 00011 1010001100

Excess-16 exp:  $00011 + 10000 = 10011$

Normalization

# Floating Point

## ■ Warming up

### ■ Example-3) 16-bit integer 1

- $1_{(10)} = \underline{0000000000000001}_{(2)}$

### ■ Example-4) 16-bit float 1.0

- $1.0_{(10)} = 1.0000000000000000_{(2)}$
- $0.100000000000000000... \times 2^1$
- (s, exp, frac) with 16-bits FPN = 0000110000000000

Excess-16 exp:  $00001 + 10000 = 10001$

# Floating Point

## ■ Fractional decimal numbers

### ■ Representation

- $d_m d_{m-1} \cdots d_1 d_0 . d_{-1} d_{-2} \cdots d_{-(n-1)} d_{-n}$

### ■ Value

- $d = \sum_{i=-n}^m 10^i \times d_i$

### ■ Example) 12.34

- $1 \times 10^1 + 2 \times 10^0 + 3 \times 10^{-1} + 4 \times 10^{-2} = 12 \frac{34}{100}$

# Floating Point

## ■ Fractional binary numbers

### ■ Representation

- $b_m b_{m-1} \cdots b_1 b_0 . b_{-1} b_{-2} \cdots b_{-(n-1)} b_{-n}$

### ■ Value

- $b = \sum_{i=-n}^m 2^i \times b_i$


### ■ Example)

- $101.11_2 = 1 \times 2^2 + 1 \times 2^0 + 1 \times 2^{-1} + 1 \times 2^{-2} = 5 \frac{3}{4}$



# Floating Point

## ■ Fractional binary numbers



# Floating Point

## ■ Examples)

Value	Representation
$5\frac{3}{4}$	$101.11_2$
$2\frac{7}{8}$	$10.111_2$
$\frac{63}{64}$	$0.111111_2$

## ■ Observations

- Divide by 2 by shifting right
- Multiply by 2 by shifting left
- Numbers of form **0.111111...<sub>2</sub>** just below **1.0**
  - $1/2 + 1/4 + 1/8 + \dots + 1/2^i + \dots \rightarrow 1.0$
  - Use notation **1.0 -  $\epsilon$**

# Floating Point

## ■ Representable numbers

- Can only exactly represent numbers of the form  $x/2^k$
- Other numbers have repeating bit representations

Value	Representation
1/3	0.0101010101[01] $\cdots_2$
1/5	0.001100110011[0011] $\cdots_2$
1/10	0.0001100110011[0011] $\cdots_2$

# Floating Point

## ■ IEEE FP representation

- Encodes rational numbers of the form  $V = x * 2^y$
- Useful for very large numbers ( $|V| \gg 0$ ) or numbers very close to 0 ( $|V| \ll 1$ )
- IEEE Standard 754 (1985-)
  - Standard for representing floating-point numbers
  - Sponsored by Intel and IEEE
  - William Kahan (1933-)
    - ✓ The father of floating point
    - ✓ Emeritus professor of EECS at the Univ. of California, Berkeley
    - ✓ Turing award, 1989



# Floating Point

## ■ IEEE FP representation

- Represents numbers in a form  $(-1)^s \times M \times 2^E$ 
  - The sign  $s$  determines whether the number is negative ( $s = 1$ ) or positive ( $s = 0$ )
    - ✓ The interpretation of the sign bit for numeric value 0 is handled as a special case
  - The significand  $M$  is a fractional binary number that ranges  $[1, 2)$
  - The exponent  $E$  weights the value by a power of 2

## ■ Encoding



- MSB is sign bit
- The **exp** encodes  $E$
- The **frac** encodes  $M$

# Floating Point

## ■ IEEE FP representation

### ▪ Encoding



- MSB is sign bit
- The **exp** encodes E
- The **frac** encodes M

### ▪ Sizes

- Single precision
  - ✓ 8 **exp** bits, 23 **frac** bits (32bits total)
- Double precision
  - ✓ 11 **exp** bits, 52 **frac** bits (64bits total)

# Floating Point


## ■ IEEE FP representation

- Encoding: 3 cases
  - Normalized values
  - Denormalized values
  - Special values

### 1. Normalized



### 2. Denormalized



### 3a. Infinity



### 3b. NaN



# Floating Point

## ■ IEEE FP representation

- Encoding: Normalized values
  - Condition
    - ✓ **exp**  $\neq$  000...0 and **exp**  $\neq$  111...1
  - Exponent coded as biased value
    - ✓ **E = Exp – Bias**  
**Exp**: unsigned value denoted by **exp**  
**Bias**: bias value ( $2^{k-1} - 1$ )
      - Single precision: 127 (**Exp**: 1~254, **E**: -126~127)
      - Double precision: 1023 (**Exp**: 1~2046, **E**: -1022~1023)
  - Significand coded with implied leading 1
    - ✓ **M = 1 + f = 1.xxx...x<sub>2</sub>** ( $0 \leq f < 1$ )
      - Minimum when 000...0 (**M** = 1.0)
      - Maximum when 111...1 (**M** = 2.0 -  $\epsilon$ )
    - ✓ Get extra leading bit for "free"



# Floating Point

## ■ IEEE FP representation

- Encoding: Normalized values (Example) **float f = 2003.0;**

- Value

$$2003_{10} = 11111010011_2 = 1.1111010011_2 \times 2^{10}$$

- Significand

$$M = 1.\underline{1111010011}_2$$

$$\text{frac} = \underline{1111010011}0000000000000000_2$$

- Exponent

$$E = 10$$

$$\text{Exp} = E + \text{Bias} = 10 + 127 = 137 = 10001001_2$$

[Floating Point Representation]

Hex	4	4	F	A	6	0	0	0
Binary	0100	0100	1111	1010	0110	0000	0000	0000
Exp	100	0100	1					
frac			1111	1010	0110			

# Floating Point

## ■ IEEE FP representation

- Encoding: Denormalized values
  - Condition
    - ✓ **exp** = 000...0
  - Exponent and significand
    - ✓ Exponent value **E** = 1 – Bias
    - ✓ Significand value **M** = 0.xxx...x<sub>2</sub> (no implied leading 1)
  - Cases
    - ✓ **exp** = 000...0, **frac** = 000...0
      - Represents value 0
      - Note that we have distinct values +0 and -0
    - ✓ **exp** = 000...0, **frac** ≠ 000...0
      - Numbers very close to 0.0
      - **Gradual underflow** property  
(possible numeric values are spaced evenly near 0.0)

# Floating Point

## ■ IEEE FP representation

### ▪ Encoding: Special values

- Condition

- ✓ **exp** = 111...1

- Cases

- ✓ **exp** = 111...1, **frac** = 000...0

- Represents value  $\pm\infty$  (infinity)
    - Represents results that overflow
    - Eg)  $1.0/0.0 = -1.0/-0.0 = +\infty$ ,  $1.0/-0.0 = -\infty$


- ✓ **exp** = 111...1, **frac**  $\neq$  000...0

- Not-a-Number (NaN)
    - Used when the result of an operation cannot be represented as a real number or infinity
    - Eg)  $\text{sqrt}(-1)$ ,  $\infty - \infty$ ,  $\infty * 0$ , ...

# Floating Point

## ■ Tiny FP example

- 8-bit FP representation
  - The sign bit is in the most significant bit
  - The next four bits are the **exp**, with a bias of 7 ( $2^3 - 1$ )
  - The last three bits are the **frac**
- Same general form as IEEE format
  - Normalized, denormalized
  - Representation of 0, NaN, infinity



# Floating Point

## ■ Tiny FP example: Values related to the exponent (*Bias* = 7)

Description	Exp	exp	E = Exp - Bias	2 <sup>E</sup>
Denormalized	0	0000	-6	1/64
Normalized	1	0001	-6	1/64
	2	0010	-5	1/32
	3	0011	-4	1/16
	4	0100	-3	1/8
	5	0101	-2	1/4
	6	0110	-1	1/2
	7	0111	0	1
	8	1000	1	2
	9	1001	2	4
	10	1010	3	8
	11	1011	4	16
	12	1100	5	32
	13	1101	6	64
	14	1110	7	128
inf, NaN	15	1111	-	-

# Floating Point

## ■ Tiny FP example: Dynamic range

Description	Bit representation	e	E	f	M	V
Zero	0 0000 000	0	-6	0	0	0
Smallest positive	0 0000 001	0	-6	1/8	1/8	1/512
	0 0000 010	0	-6	2/8	2/8	2/512
	0 0000 011	0	-6	3/8	3/8	3/512
	0 0000 110	0	-6	6/8	6/8	6/512
Largest denorm.	0 0000 111	0	-6	7/8	7/8	7/512
Smallest norm.	0 0001 000	1	-6	0	8/8	8/512
	0 0001 001	1	-6	1/8	9/8	9/512
	0 0110 110	6	-1	6/8	14/8	14/16
	0 0110 111	6	-1	7/8	15/8	15/16
One	0 0111 000	7	0	0	8/8	1
	0 0111 001	7	0	1/8	9/8	9/8
	0 0111 010	7	0	2/8	10/8	10/8
	0 1110 110	14	7	6/8	14/8	224
Largest norm.	0 1110 111	14	7	7/8	15/8	240
Infinity	0 1111 000	-	-	-	-	$+\infty$

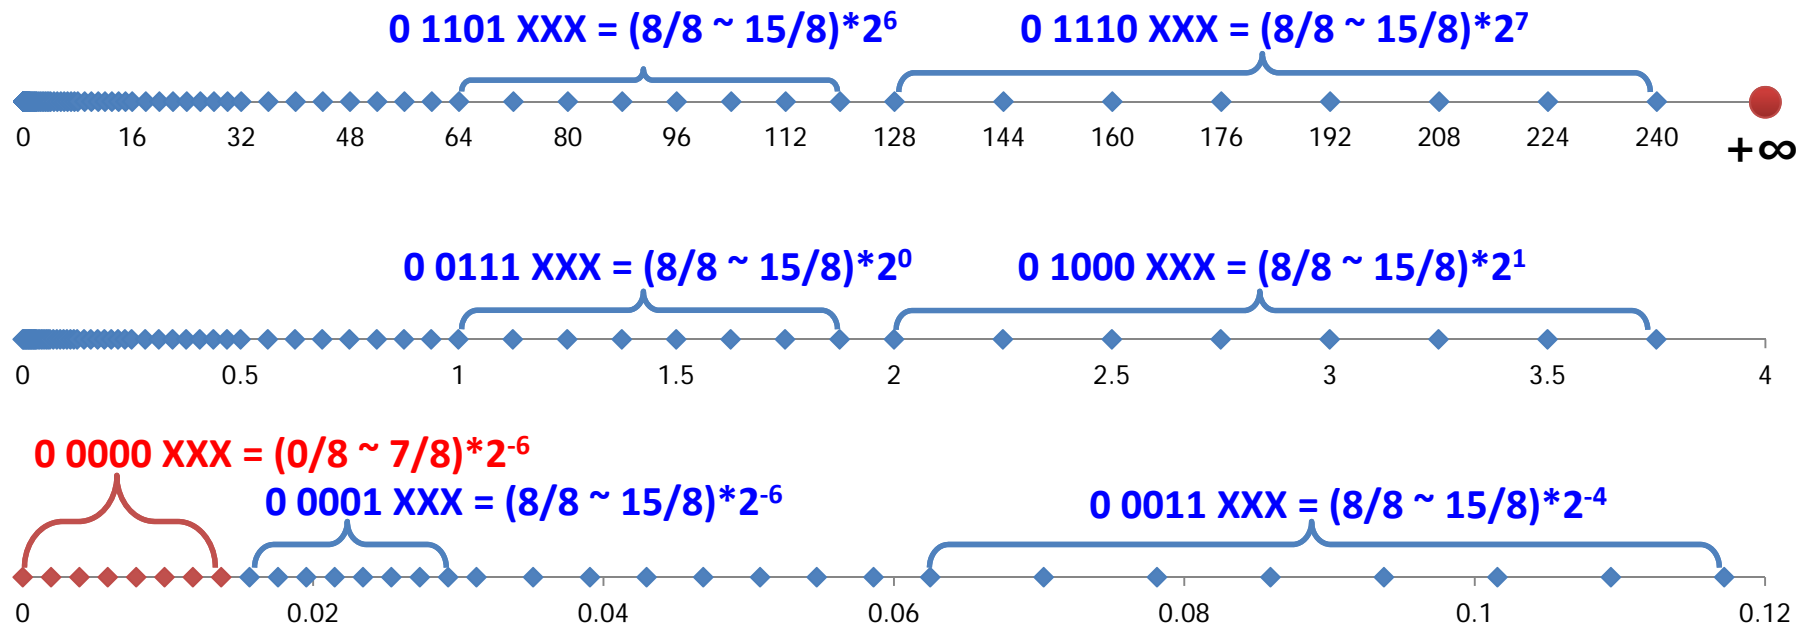
# Floating Point

## ■ Tiny FP example


Description	Bit representation	Exponent			Fraction		Value		
		$e$	$E$	$2^E$	$f$	$M$	$2^E \times M$	$V$	Decimal
Zero	0 0000 000	0	-6	$\frac{1}{64}$	0	0	$\frac{0}{512}$	0	0.0
Smallest pos.	0 0000 001	0	-6	$\frac{1}{64}$	$\frac{1}{8}$	$\frac{1}{8}$	$\frac{1}{512}$	$\frac{1}{512}$	0.001953
	0 0000 010	0	-6	$\frac{1}{64}$	$\frac{2}{8}$	$\frac{2}{8}$	$\frac{2}{512}$	$\frac{1}{256}$	0.003906
	0 0000 011	0	-6	$\frac{1}{64}$	$\frac{3}{8}$	$\frac{3}{8}$	$\frac{3}{512}$	$\frac{3}{512}$	0.005859
	⋮								
Largest denorm.	0 0000 111	0	-6	$\frac{1}{64}$	$\frac{7}{8}$	$\frac{7}{8}$	$\frac{7}{512}$	$\frac{7}{512}$	0.013672
Smallest norm.	0 0001 000	1	-6	$\frac{1}{64}$	0	$\frac{8}{8}$	$\frac{8}{512}$	$\frac{1}{64}$	0.015625
	0 0001 001	1	-6	$\frac{1}{64}$	$\frac{1}{8}$	$\frac{9}{8}$	$\frac{9}{512}$	$\frac{9}{512}$	0.017578
	⋮								
	0 0110 110	6	-1	$\frac{1}{2}$	$\frac{6}{8}$	$\frac{14}{8}$	$\frac{14}{16}$	$\frac{7}{8}$	0.875
One	0 0110 111	6	-1	$\frac{1}{2}$	$\frac{7}{8}$	$\frac{15}{8}$	$\frac{15}{16}$	$\frac{15}{16}$	0.9375
	0 0111 000	7	0	1	0	$\frac{8}{8}$	$\frac{8}{8}$	1	1.0
	0 0111 001	7	0	1	$\frac{1}{8}$	$\frac{9}{8}$	$\frac{9}{8}$	$\frac{9}{8}$	1.125
	0 0111 010	7	0	1	$\frac{2}{8}$	$\frac{10}{8}$	$\frac{10}{8}$	$\frac{5}{4}$	1.25
	⋮								
	0 1110 110	14	7	128	$\frac{6}{8}$	$\frac{14}{8}$	$\frac{1792}{8}$	224	224.0
Largest norm.	0 1110 111	14	7	128	$\frac{7}{8}$	$\frac{15}{8}$	$\frac{1920}{8}$	240	240.0
Infinity	0 1111 000	—	—	—	—	—	—	$\infty$	—

# Floating Point

## ■ Tiny FP example: Encoded values (nonnegative)



(Without denormalization)







# Floating Point

## ■ Tiny FP example: Encoded values for 6-bit FP

- 3 exponent bits and 2 fraction bits (bias = 3)



# Floating Point

## ■ Interesting numbers

Description	exp	frac	Numeric value
Zero	000 ... 00	000 ... 00	0.0
Smallest Positive Denormalized	000 ... 00	000 ... 01	Single: $2^{-23} \times 2^{-126} \approx 1.4 \times 10^{-45}$ Double: $2^{-52} \times 2^{-1022} \approx 4.9 \times 10^{-324}$
Largest Denormalized	000 ... 00	111 ... 11	Single: $(1.0 - \epsilon) \times 2^{-126} \approx 1.18 \times 10^{-38}$ Double: $(1.0 - \epsilon) \times 2^{-1022} \approx 2.2 \times 10^{-308}$
Smallest Positive Normalized	000 ... 01	000 ... 00	Single: $1.0 \times 2^{-126}$ , Double: $1.0 \times 2^{-1022}$ (Just larger than largest denormalized)
One	011 ... 11	000 ... 00	1.0
Largest Normalized	111 ... 10	111 ... 11	Single: $(2.0 - \epsilon) \times 2^{127} \approx 3.4 \times 10^{38}$ Double: $(2.0 - \epsilon) \times 2^{1023} \approx 1.8 \times 10^{308}$

# Floating Point

## ■ Special properties

- FP zero same as integer zero
  - All bits = 0
  
- To compare two FP numbers
  - Can (almost) use unsigned integer comparison with some considerations
    - ✓ Must first compare sign bits
    - ✓ Must consider  $-0 = +0$
    - ✓ NaNs problematic
      - Will be greater than any other values
    - ✓ Otherwise OK
      - Denormalized vs normalized
      - Normalized vs infinity

# Floating Point

## ■ Rounding

- A systematic method of finding the “closest” matching value  $x'$  that can be represented in the desired FP format, for a given value  $x$
- IEEE FP format defines 4 different rounding modes
  - Round-to-even (round-to-nearest)
  - Round-toward-zero
  - Round-down
  - Round-up

# Floating Point

## ■ Rounding

- Round-to-even mode (default)
  - Attempts to find a closest match
  - For halfway values
    - ✓ Rounds upward or downward such that the LSB of the result is even
    - ✓ Avoids the statistical bias (for example, in averaging)
- Round-toward-zero
  - Rounds positive numbers downward and negative numbers upward
  - Takes the value  $\mathbf{x}^*$  such that  $|\mathbf{x}^*| \leq |\mathbf{x}|$
- Round-down
  - Always rounds downward (takes the value  $\mathbf{x}^-$  such that  $\mathbf{x}^- \leq \mathbf{x}$ )
- Round-up
  - Always rounds upward (takes the value  $\mathbf{x}^+$  such that  $\mathbf{x}^+ \geq \mathbf{x}$ )

# Floating Point



## ■ Rounding

### ■ Example)

Mode	\$1.40	\$1.60	\$1.50	\$2.50	\$-1.50
Round-to-even	\$1	\$2	\$2	\$2	\$-2
Round-toward-zero	\$1	\$1	\$1	\$2	\$-1
Round-down	\$1	\$1	\$1	\$2	\$-2
Round-up	\$2	\$2	\$2	\$3	\$-1

# Floating Point

## ■ Rounding: Round-to-even on binary fractional numbers

### ■ Consider

- LSB 0: even
- LSB 1: odd

### ■ Example) Rounding to the nearest quarter

- $10.00011_{(2)} \rightarrow 10.00_{(2)} \quad 2\frac{3}{32} \rightarrow 2$
- $10.00110_{(2)} \rightarrow 10.01_{(2)} \quad 2\frac{3}{16} \rightarrow 2\frac{1}{4}$
- $10.11100_{(2)} \rightarrow 11.00_{(2)} \quad 2\frac{7}{8} \rightarrow 3$
- $10.10100_{(2)} \rightarrow 10.10_{(2)} \quad 2\frac{5}{8} \rightarrow 2\frac{1}{2}$

# Floating Point

## ■ FP operations: Multiplication

- Operands
$$(-1)^{s1} \mathbf{M1} 2^{E1} \times (-1)^{s2} \mathbf{M2} 2^{E2}$$
- Exact Result
$$(-1)^s \mathbf{M} 2^E$$
  - Sign **s**:  $s1 \wedge s2$
  - Significand **M**:  $\mathbf{M1} * \mathbf{M2}$
  - Exponent **E**:  $E1 + E2$
- Fixing
  - If  $\mathbf{M} \geq 2$ , shift **M** right, increment **E**
  - If **E** out of range, overflow
  - Round **M** to fit **frac** precision



# Floating Point

## ■ FP operations: Addition

### ■ Operands

$$(-1)^{s1} \mathbf{M1} 2^{E1} + (-1)^{s2} \mathbf{M2} 2^{E2} \quad (\text{Assume } E1 > E2)$$

### ■ Exact Result

$$(-1)^s \mathbf{M} 2^E$$

- Sign **s**, significand **M**: result of signed align & add
- Exponent **E**: **E1**

### ■ Fixing

- If **M**  $\geq 2$ , shift **M** right, increment **E**
- If **M**  $< 1$ , shift **M** left k positions, decrement **E** by k
- Overflow if **E** out of range
- Round **M** to fit **frac** precision

# Floating Point

## ■ FP operations: Some properties

- FP addition is not associative
  - $(3.14 + 1e10) - 1e10 \rightarrow 0.0$  (3.14 is lost due to rounding)
  - $3.14 + (1e10 - 1e10) \rightarrow 3.14$
- FP addition satisfies monotonicity
  - If  $a \geq b$ , then  $x + a \geq x + b$ , for any  $a, b, x$  other than NaN
- FP multiplication is not associative
  - $(1e20 * 1e20) * 1e-20 \rightarrow \infty$  (in single precision arithmetic)
  - $1e20 * (1e20 * 1e-20) \rightarrow 1e20$
- FP multiplication does not distribute over addition
  - $1e20 * (1e20 - 1e20) \rightarrow 0.0$
  - $1e20 * 1e20 - 1e20 * 1e20 \rightarrow \text{NaN}$

# Floating Point

## ■ FP in C

- C standard provides two FP types
  - **float** and **double**
- C standard does not require the machine to use IEEE FP
  - Also, no standard methods to change the rounding mode or to get special values such as  $-0$ ,  $+\infty$ ,  $-\infty$ , or NaN
- Most systems provide the header files and libraries to provide access to these features
  - But, the details vary from one system to another

# Floating Point

## ■ FP in C

- 3<sup>rd</sup> FP types in ISO C99: **long double**
  - Equivalent to **double** type in many machines
  - 80-bit extended precision format in Intel-compatible machines
    - ✓ 15 **exp** bits and 63 **frac** bits (1 bit wasted)

# Floating Point

## ■ FP in C

- Type conversions
  - Casting between **int**, **float**, and **double** changes numeric values
  - From **int** to **float**
    - ✓ May be rounded
  - From **int** or **float** to **double**
    - ✓ Exact conversion
  - From **double** to **float**
    - ✓ May overflow or may be rounded
  - From **float** or **double** to **int**
    - ✓ Rounded toward 0
    - ✓ May overflow
      - No specifications in C standard

# Floating Point

## ■ FP puzzles

```
int x, float f, double d; //neither d nor f is NaN
```

- `x == (int)(float) x;`
- `x == (int)(double) x;`
- `f == (float)(double) f;`
- `d == (float) d;`
- `f == -(-f);`
- `2/3 == 2/3.0;`
- `d < 0.0  $\Rightarrow$  ((d*2) < 0.0);`
- `d > f  $\Rightarrow$  -f > -d;`
- `d * d >= 0.0;`
- `(d + f) - d == f;`

# Floating Point

## ■ Ariane 5 tragedy (Jun. 4, 1996)

- Exploded 37 seconds after liftoff
- Satellites worth \$500 million

## ■ Why?

- Computed horizontal velocity as FP number
- Converted to 16-bit integer
  - Careful analysis of **Ariane 4** trajectory proved 16-bit is enough
- Reused a module from 10-year-old SW
  - Overflowed for **Ariane 5**
  - No precise specification for the SW



# Summary

## ■ IEEE FP has clear mathematical properties

- Represents numbers of form  $M \times 2^E$
- Can reason about operations independent of implementation
  - As if computed with perfect precision and then rounded
- Not the same as real arithmetic
  - Lacks associativity/distributivity
  - Makes life difficult  
for compilers and serious numerical applications programmers