[Chap.1] A Tour on Computer Systems

Young Ik Eom (yieom@skku.edu, 031-290-7120)
Distributing Computing Laboratory
Sungkyunkwan University
http://dclab.skku.ac.kr

SKK UNIVERSITY

Goal

Understand what happens and why

- When you run hello program on your system

```
#include <stdio.h>

int main()
{
   printf("hello, world\n");
   return 0;
}
```

Source File

Source file

```
#include <stdio.h>
int main()
{
   printf("hello, world\n");
   return 0;
}
```

directive

#	i	n	c	1	u	d	e	SP	<	s	t	d	i	o	.
35	105	110	99	108	117	100	101	32	60	115	116	100	105	111	46
h	>	\n	\n	i	n	t	SP	m	a	i	n	(	)	\n	{
104	62	10	10	105	110	116	32	109	97	105	110	40	41	10	123
\n	SP	SP	SP	SP	p	r	i	n	t	f	(	"	h	e	l
10	32	32	32	32	112	114	105	110	116	102	40	34	104	101	108
l	o	,	SP	w	o	r	l	d	\	n	"	)	;	\n	SP
108	111	44	32	119	111	114	108	100	92	110	34	41	59	10	32
SP	SP	SP	r	e	t	u	r	n	SP	0	;	\n	}	\n	
32	32	32	114	101	116	117	114	110	32	48	59	10	125	10	

Source File

Notes)

- Text file
  - File that consists exclusively of ASCII characters
- Binary file
  - …
- ASCII standard
  - Represents each character with a unique 8-bit integer value

Source File (ASCII Code Table)

Dec	Hx	Oct	Char	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr
0	0	000	NUL (null)	32	20	040		Space	64	40	100	@	@	96	60	140	`	`
1	1	001	SOH (start of heading)	33	21	041	!	!	65	41	101	A	A	97	61	141	a	a
2	2	002	STX (start of text)	34	22	042	"	"	66	42	102	B	B	98	62	142	b	b
3	3	003	ETX (end of text)	35	23	043	#	#	67	43	103	C	C	99	63	143	c	c
4	4	004	EOT (end of transmission)	36	24	044	$	$	68	44	104	D	D	100	64	144	d	d
5	5	005	ENQ (enquiry)	37	25	045	%	%	69	45	105	E	E	101	65	145	e	e
6	6	006	ACK (acknowledge)	38	26	046	&	&	70	46	106	F	F	102	66	146	f	f
7	7	007	BEL (bell)	39	27	047	'	'	71	47	107	G	G	103	67	147	g	g
8	8	010	BS (backspace)	40	28	050	(	(	72	48	110	H	H	104	68	150	h	h
9	9	011	TAB (horizontal tab)	41	29	051	)	)	73	49	111	I	I	105	69	151	i	i
10	A	012	LF (NL line feed, new line)	42	2A	052	*	*	74	4A	112	J	J	106	6A	152	j	j
11	B	013	VT (vertical tab)	43	2B	053	+	+	75	4B	113	K	K	107	6B	153	k	k
12	C	014	FF (NP form feed, new page)	44	2C	054	,	,	76	4C	114	L	L	108	6C	154	l	l
13	D	015	CR (carriage return)	45	2D	055	-	-	77	4D	115	M	M	109	6D	155	m	m
14	E	016	SO (shift out)	46	2E	056	.	.	78	4E	116	N	N	110	6E	156	n	n
15	F	017	SI (shift in)	47	2F	057	/	/	79	4F	117	O	O	111	6F	157	o	o
16	10	020	DLE (data link escape)	48	30	060	0	0	80	50	120	P	P	112	70	160	p	p
17	11	021	DC1 (device control 1)	49	31	061	1	1	81	51	121	Q	Q	113	71	161	q	q
18	12	022	DC2 (device control 2)	50	32	062	2	2	82	52	122	R	R	114	72	162	r	r
19	13	023	DC3 (device control 3)	51	33	063	3	3	83	53	123	S	S	115	73	163	s	s
20	14	024	DC4 (device control 4)	52	34	064	4	4	84	54	124	T	T	116	74	164	t	t
21	15	025	NAK (negative acknowledge)	53	35	065	5	5	85	55	125	U	U	117	75	165	u	u
22	16	026	SYN (synchronous idle)	54	36	066	6	6	86	56	126	V	V	118	76	166	v	v
23	17	027	ETB (end of trans. block)	55	37	067	7	7	87	57	127	W	W	119	77	167	w	w
24	18	030	CAN (cancel)	56	38	070	8	8	88	58	130	X	X	120	78	170	x	x
25	19	031	EM (end of medium)	57	39	071	9	9	89	59	131	Y	Y	121	79	171	y	y
26	1A	032	SUB (substitute)	58	3A	072	:	:	90	5A	132	Z	Z	122	7A	172	z	z
27	1B	033	ESC (escape)	59	3B	073	;	;	91	5B	133	[	[	123	7B	173	{	{
28	1C	034	FS (file separator)	60	3C	074	<	<	92	5C	134	\	\	124	7C	174			|
29	1D	035	GS (group separator)	61	3D	075	=	=	93	5D	135	]	]	125	7D	175	}	}
30	1E	036	RS (record separator)	62	3E	076	>	>	94	5E	136	^	^	126	7E	176	~	~
31	1F	037	US (unit separator)	63	3F	077	?	?	95	5F	137	_	_	127	7F	177	&#127;	DEL

Source: www.LookupTables.com

Translation

Compilation system

```
unix> gcc -o hello hello.c
```

4 phases

- Preprocessing
- Compilation
- Assembly
- Linking

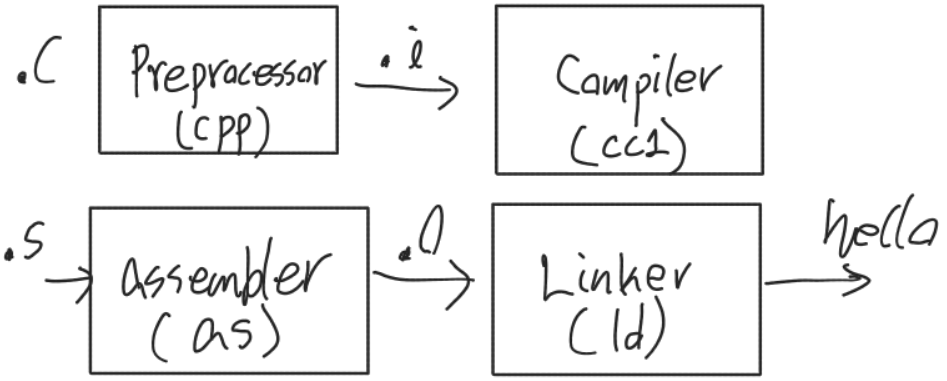

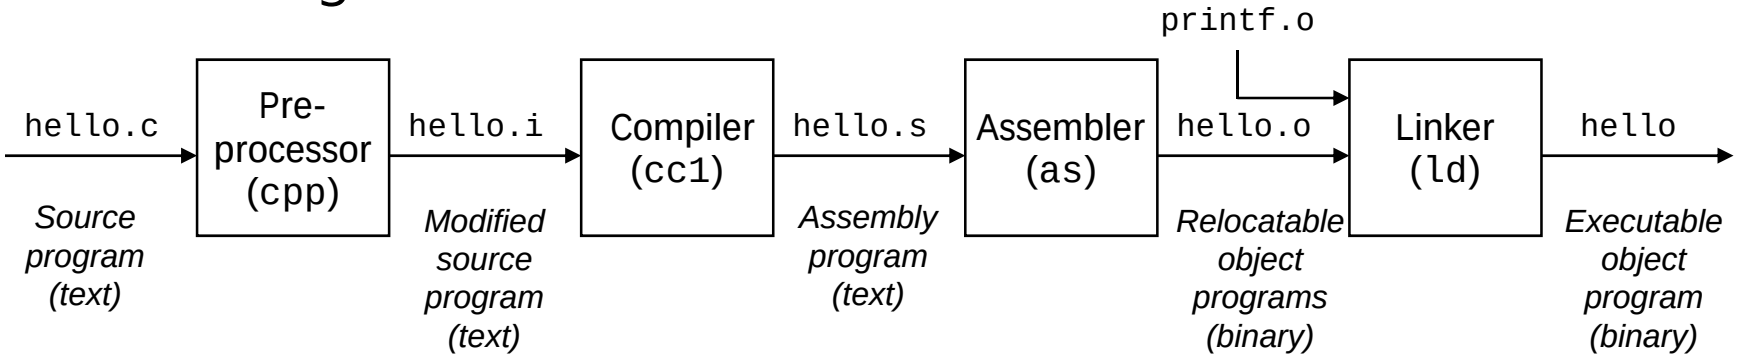

Translation

Preprocessing

- Produces pure C-language program
  - By processing preprocessor directives such as…

```
#include …
#define …
#ifdef …
```

etc…

Translation

Compiling the source program

- Produces assembly-language program

```
1   main:
2     subq   $8, %rsp
3     movl   $.LC0, %edi
4     call   puts
5     movl   $0, %eax
6     addq   $8, %rsp
7     ret
```

Translation

Relocatable object

Executable object

- **Assembling and Linking**
  - Produces binary executable code

Running

Running the executable

```
unix> ./hello
hello, world
unix> _
```

- The shell interprets the command line
  - Checks the first word in the command line corresponds to built-in shell command
  - When it is, runs the command by itself
  - When it is not, loads the executable file and runs it
  - Waits for it to terminate
  - Prompts again

Running (HW organization)

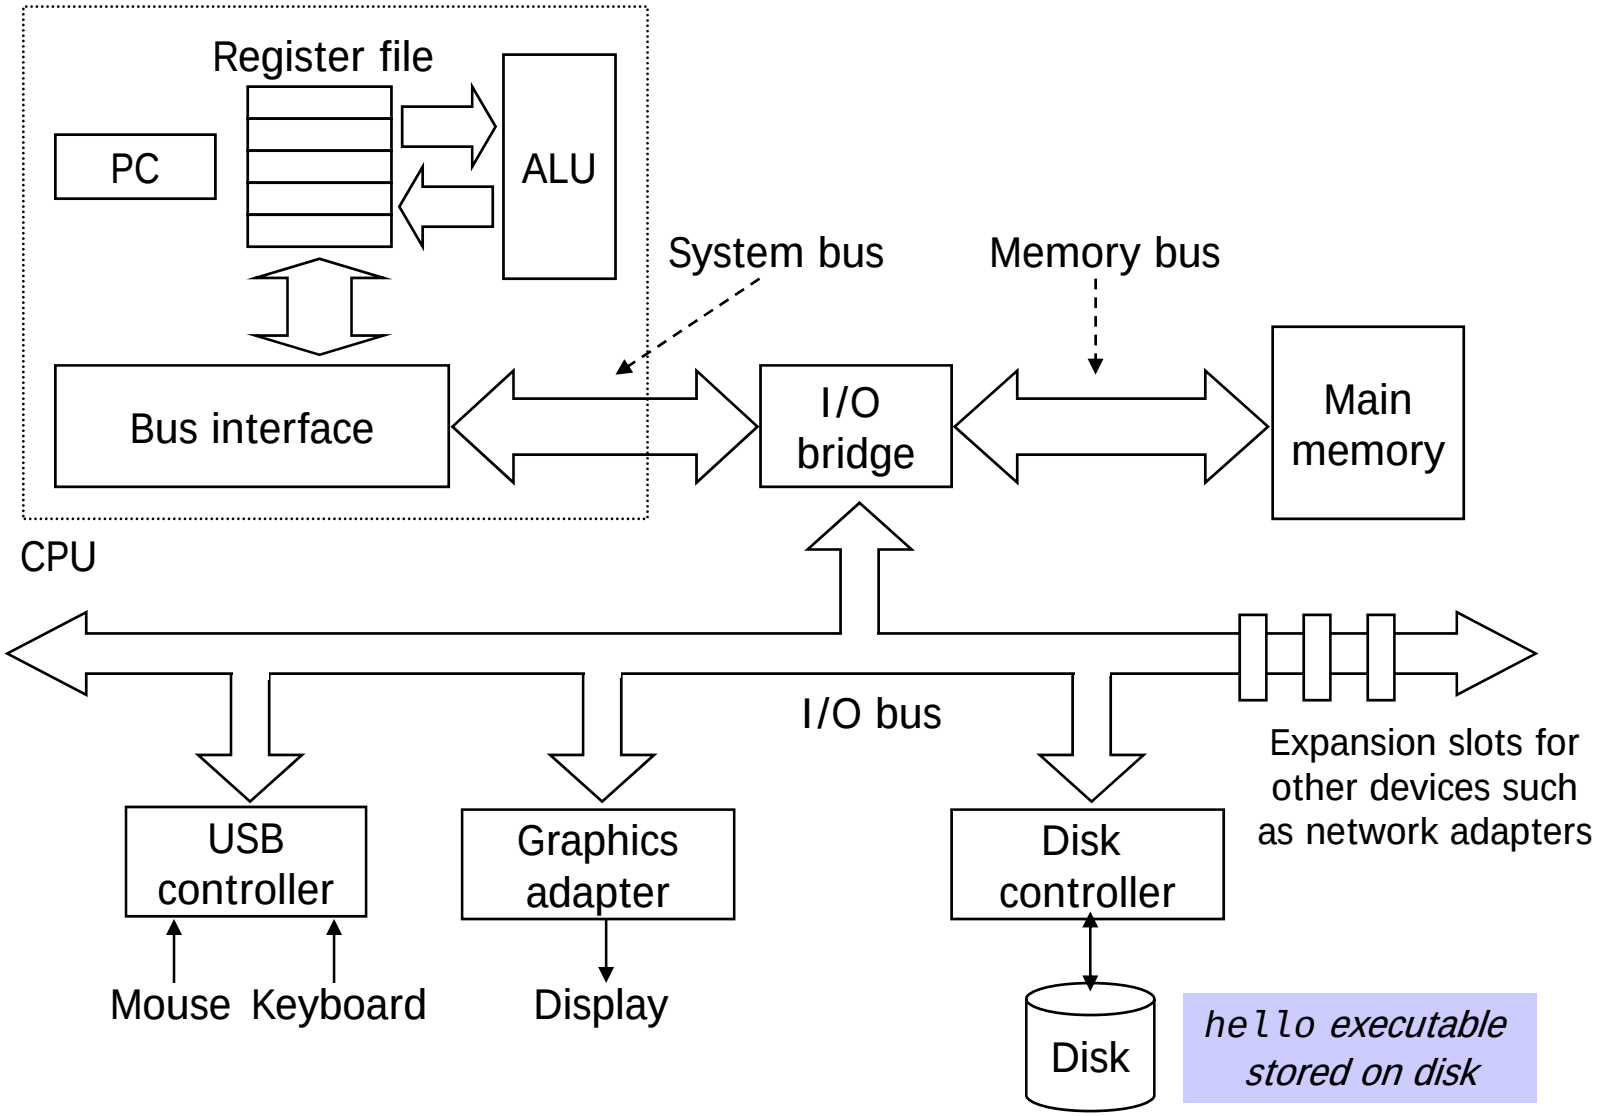

Running (HW organization)

Notes)

- Bus and word
  - Word size: 1/2/4/8 bytes
- Controller and adapter
  - Controller
    - ✓ Chipset in the device itself or on the system's main PCB
  - Adapter
    - ✓ Card that plugs into a slot on the PCB
- Instruction cycle
  - IF/ID/OF/EX/IC

Running

■ Storing the command into memory

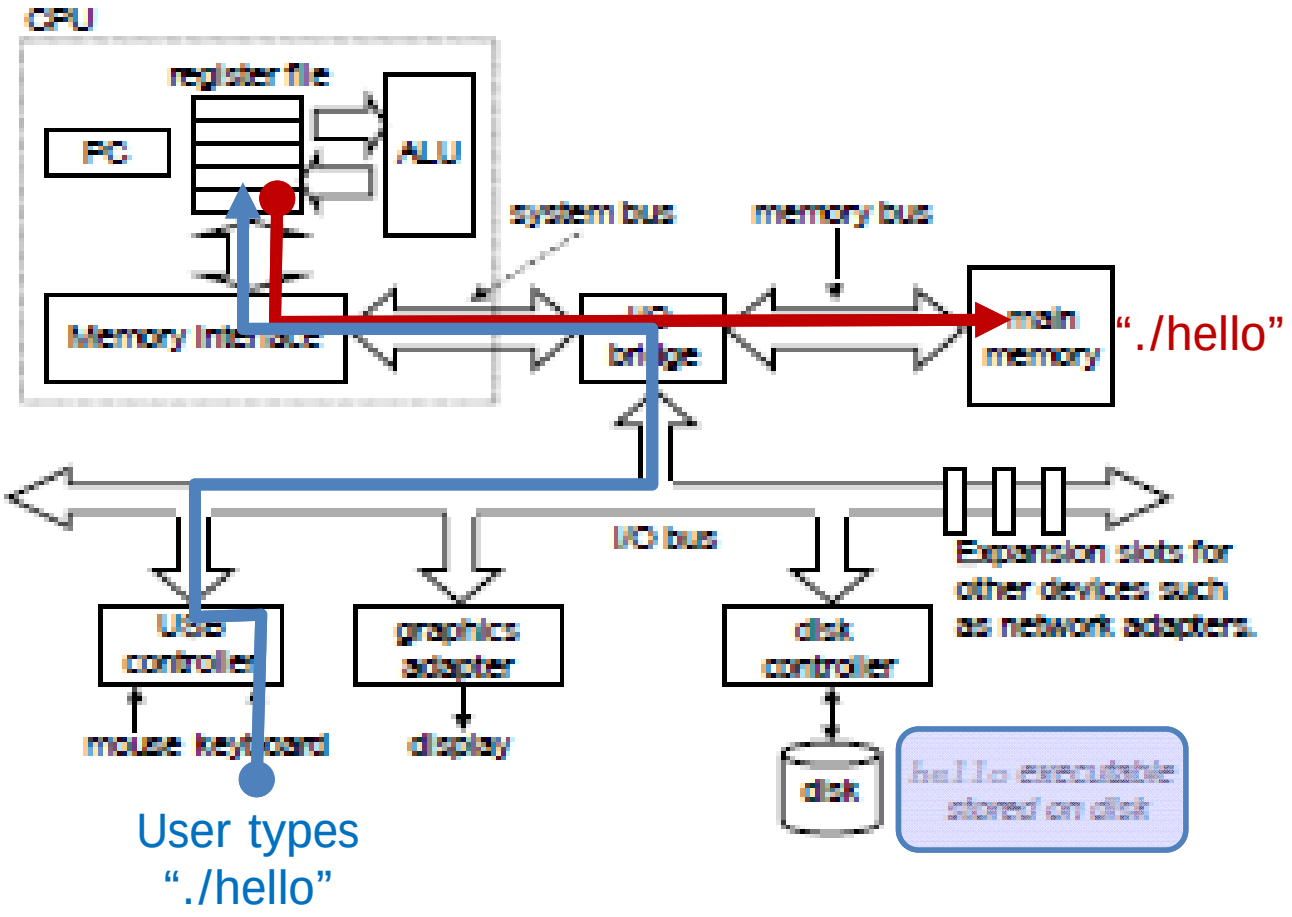

Running

Loading the executable

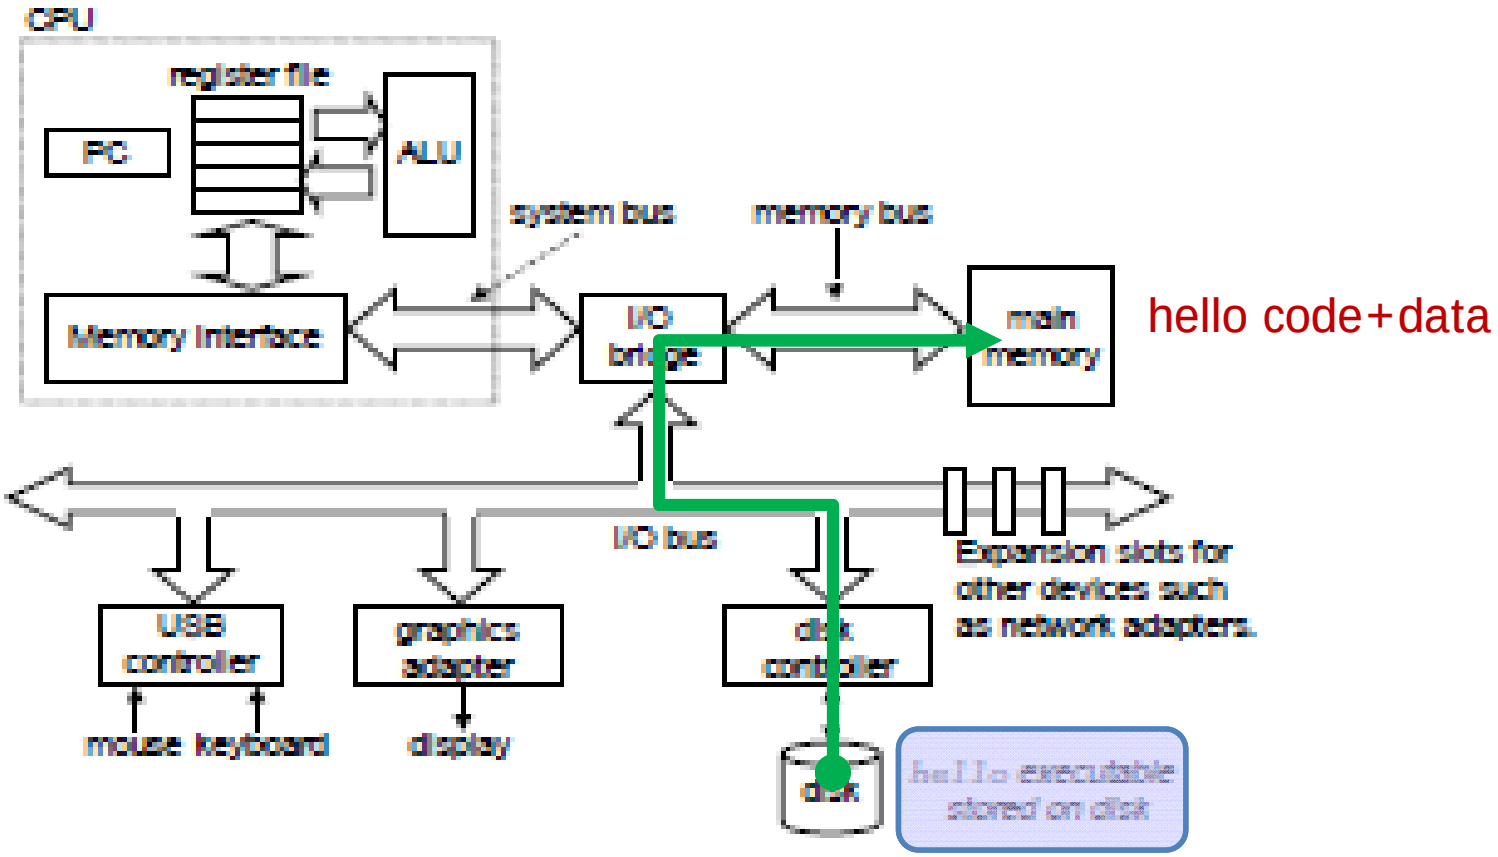

Running

Executing the program

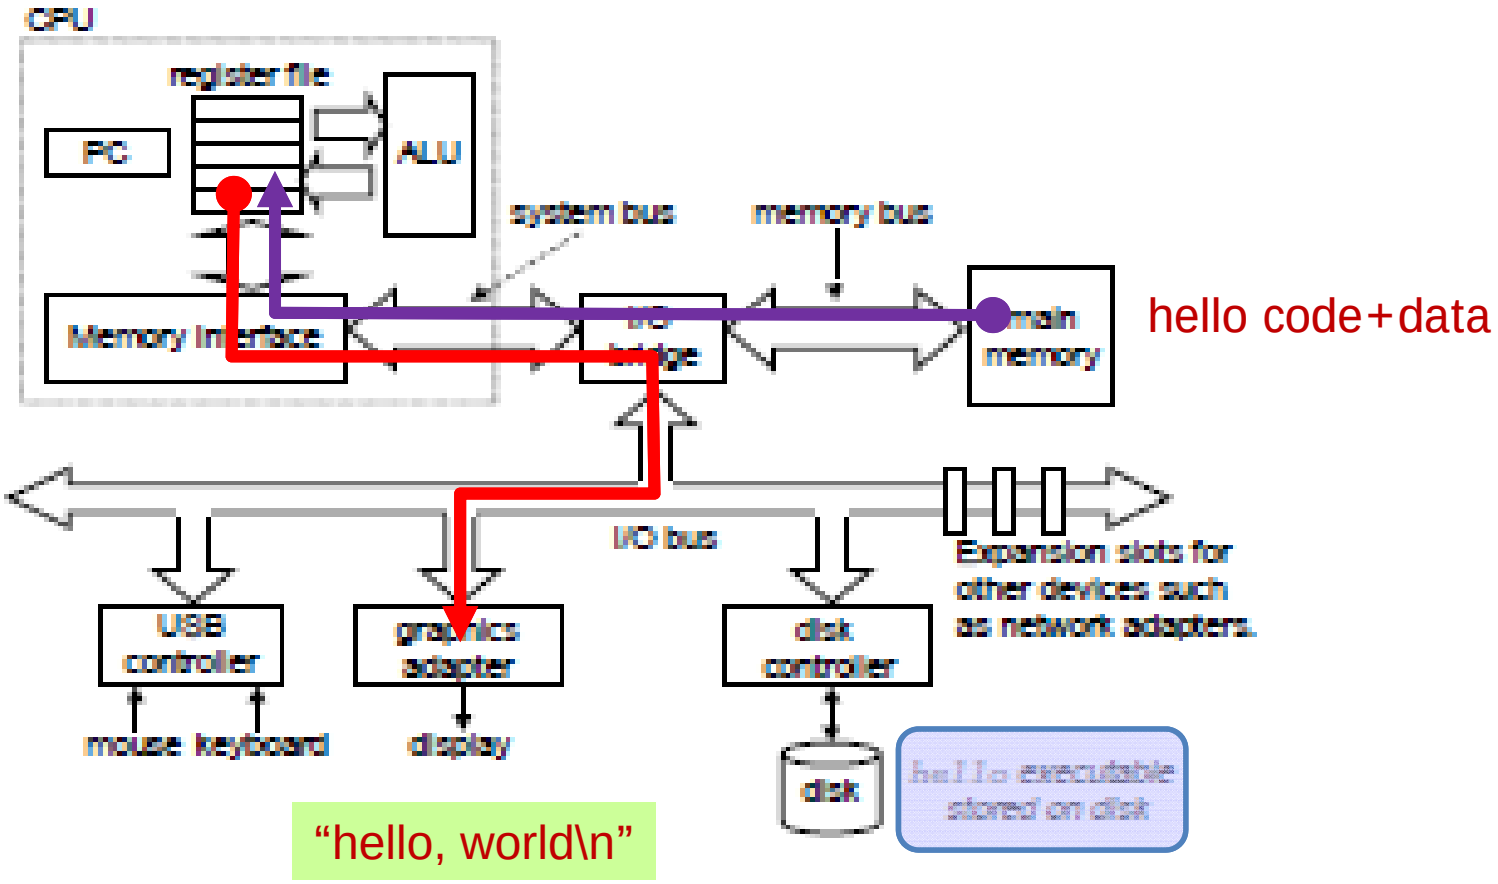

Caches

Cache memories

- Deals with the processor-memory gap
  - L1 cache (on the processor chip) 4 cycles
  - L2 cache (connected to the processor by a special bus) 10 cycles
  - L3 cache 50 cycles
- Exploits locality

한 번 접근된 명령, 데이터는 재접근될 확률↑ + 그 근처의 데이터도 접근할 확률↑

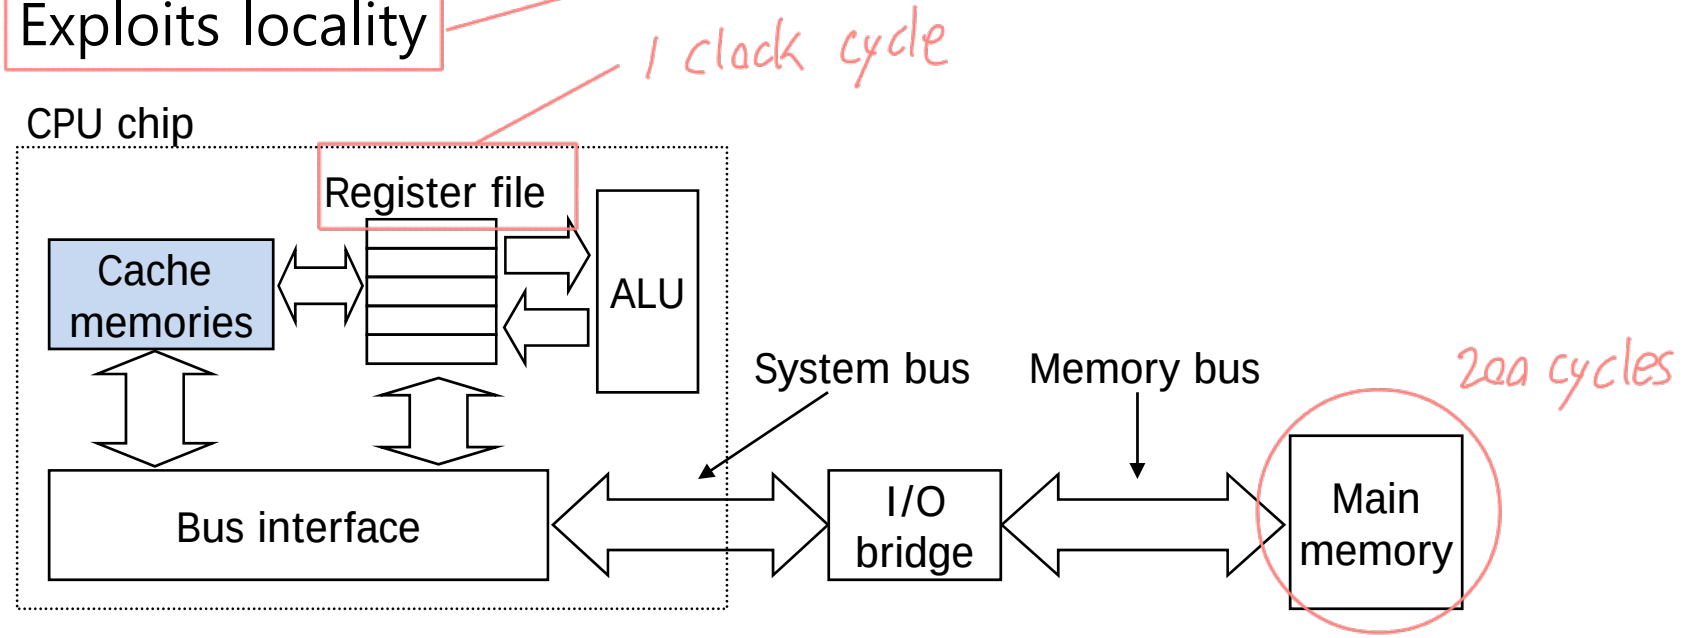

Multi-core Architectures

Multi-core processor

- Have several processors (cores) integrated into a single integrated-circuit chip

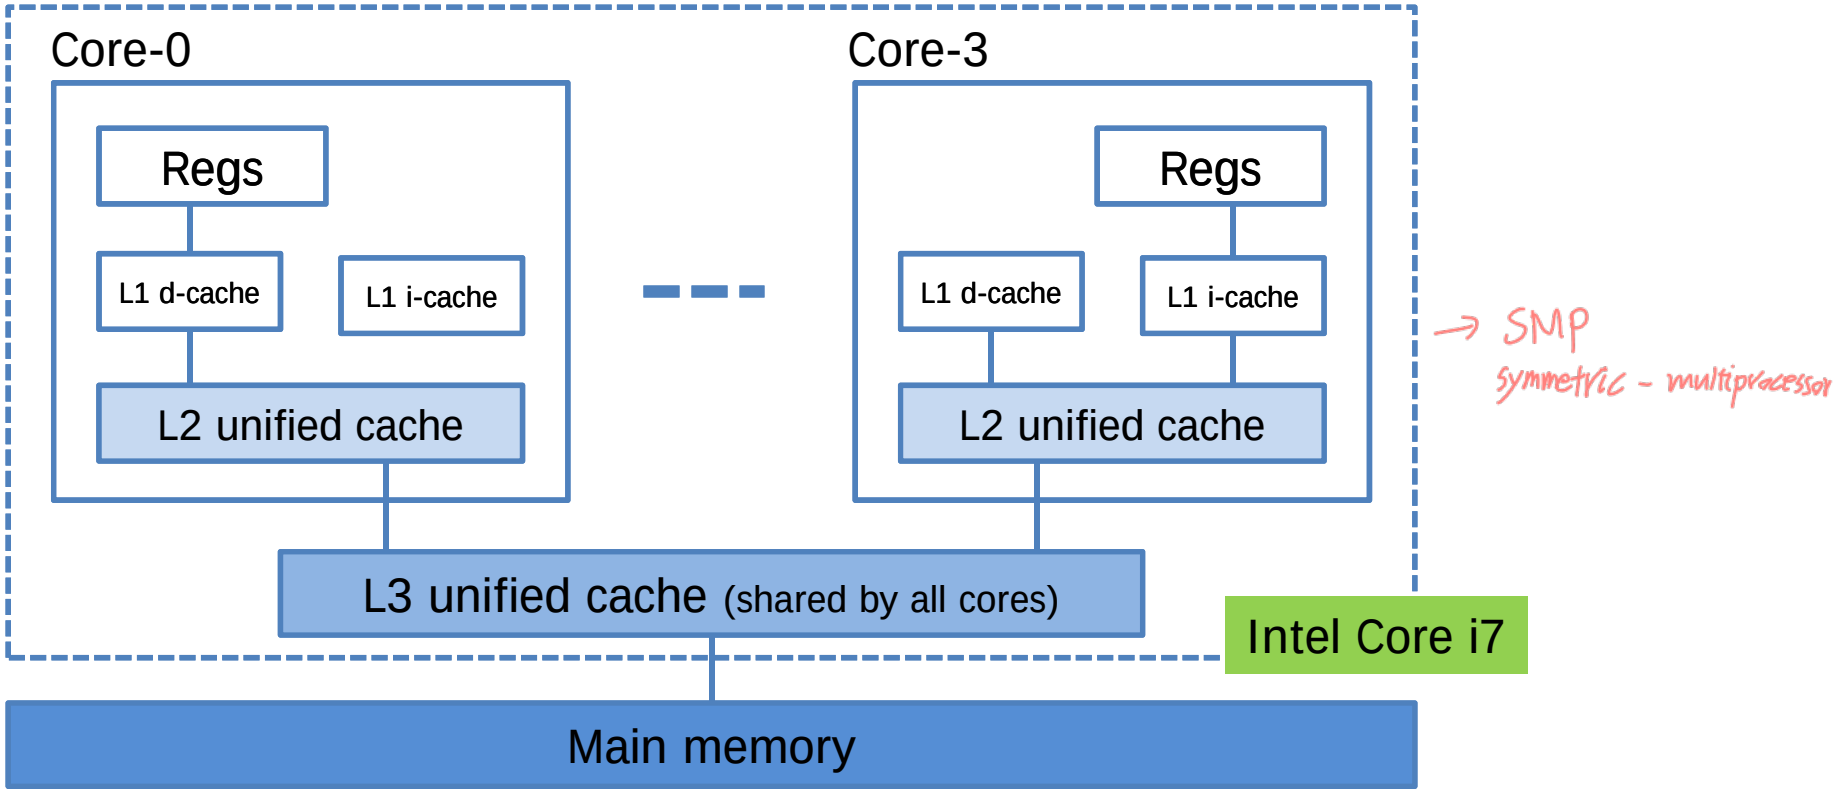

Storage Hierarchy

Hierarchy of storage devices

- Storage at one level serves as a cache for storage at the next lower level

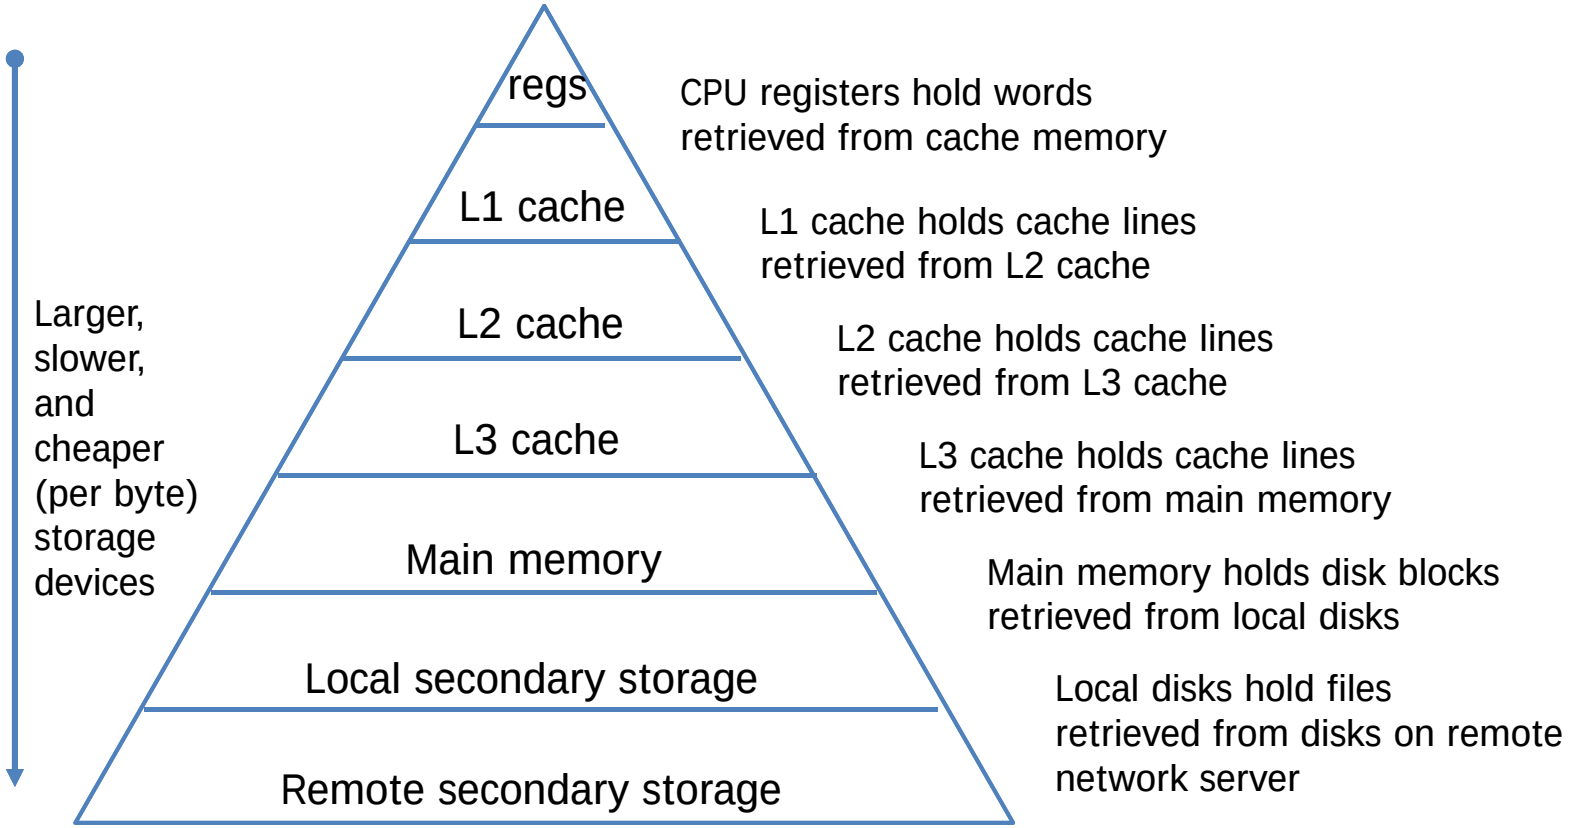

Storage Hierarchy

Storage hierarchy

Type	What cached	Where cached	Latency (#cycles)	Managed by
CPU registers	4B or 8B words	On-chip CPU registers	0~1	Hardware
L1 cache	64B blocks	On-chip L1 cache	4	Hardware
L2 cache	64B blocks	On-chip L2 cache	10	Hardware
L3 cache	64B blocks	On-chip L3 cache	50	Hardware
Main memory	4KB pages	Main memory	200	Hardware + OS
Storage (HDD)	Parts of files	Local HDD	10,000,000	OS
Web cache	Web pages	Remote server disks	1,000,000,000	Web proxy server

Operating Systems

OS

- Functions of OS
  - User interface (for user convenience)
  - Resource management (for efficiency)
  - Process management
  - + Networking / Security&Protection

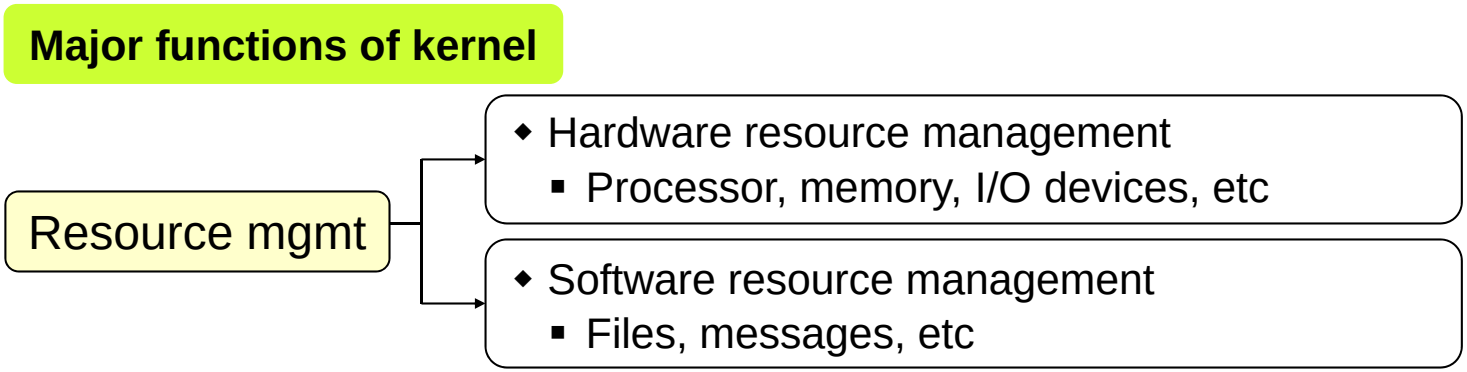

Operating Systems

OS

- Provide applications with simple and uniform interface for manipulating complicated and often widely different low-level HW devices
- Protects the HW from misuse by applications

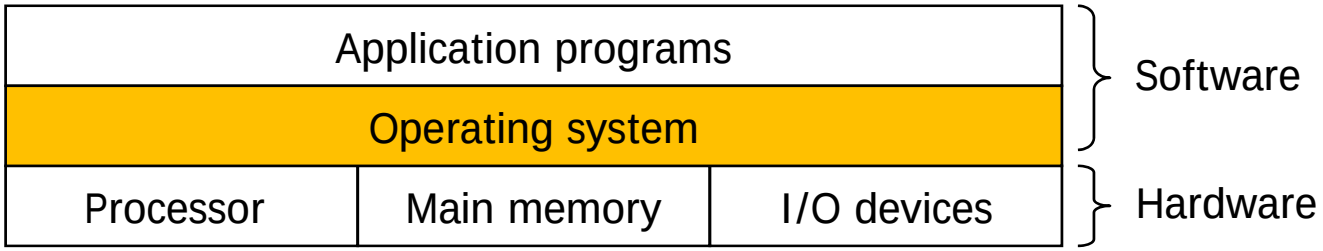

Operating Systems

OS abstractions

- Processes
- Virtual memory
- Files

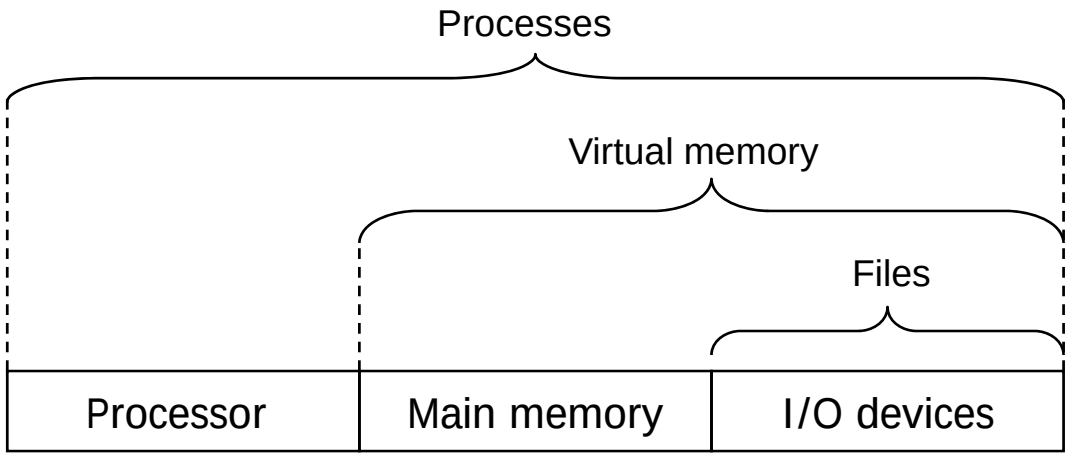

OS abstractions

- Process
  - OS's abstraction for a running program

- Context switching

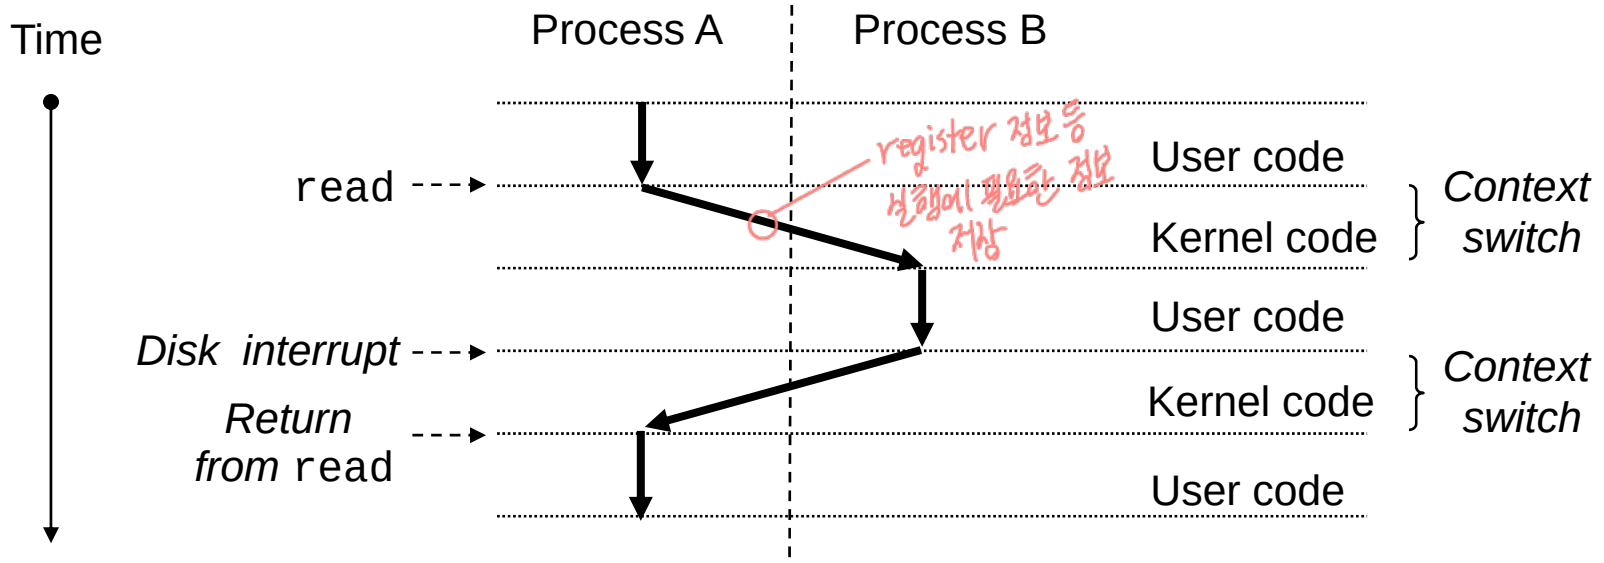

Operating Systems

OS abstractions

- Virtual memory
  - Provides each process with the illusion that it has exclusive use of main memory
- Virtual address space
  - View of memory by each process
    - ✓ Program code and data
    - ✓ Heap
    - ✓ Shared libraries
    - ✓ Stack
    - ✓ Kernel space

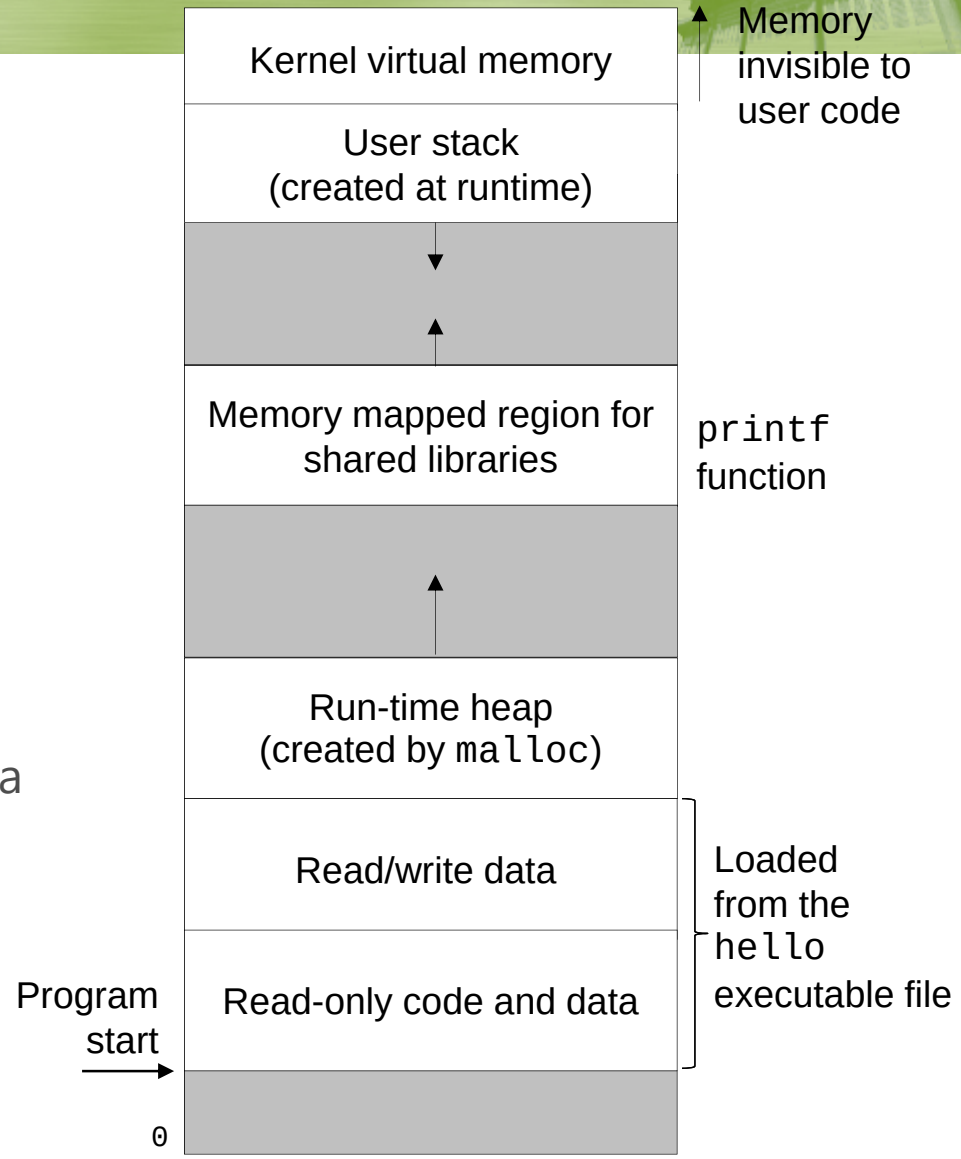

Operating Systems

OS abstractions

- File
  - Sequence of bytes
  - Provides applications with a uniform view of all of the varied IO devices that might be contained in the system

Operating Systems

More abstractions: virtual machine

- An abstraction of the entire system, including operating system, the processor, and the programs (processes)

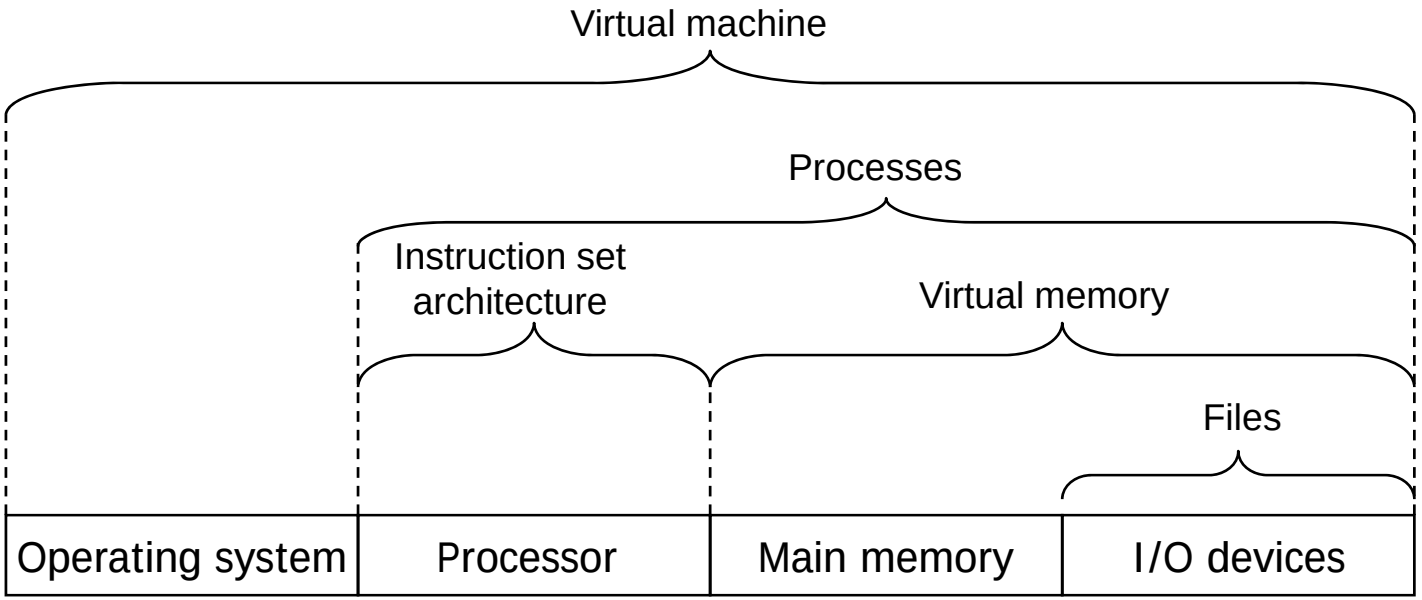

Networks

Network

- Can be viewed as just another IO device, in the view of individual system

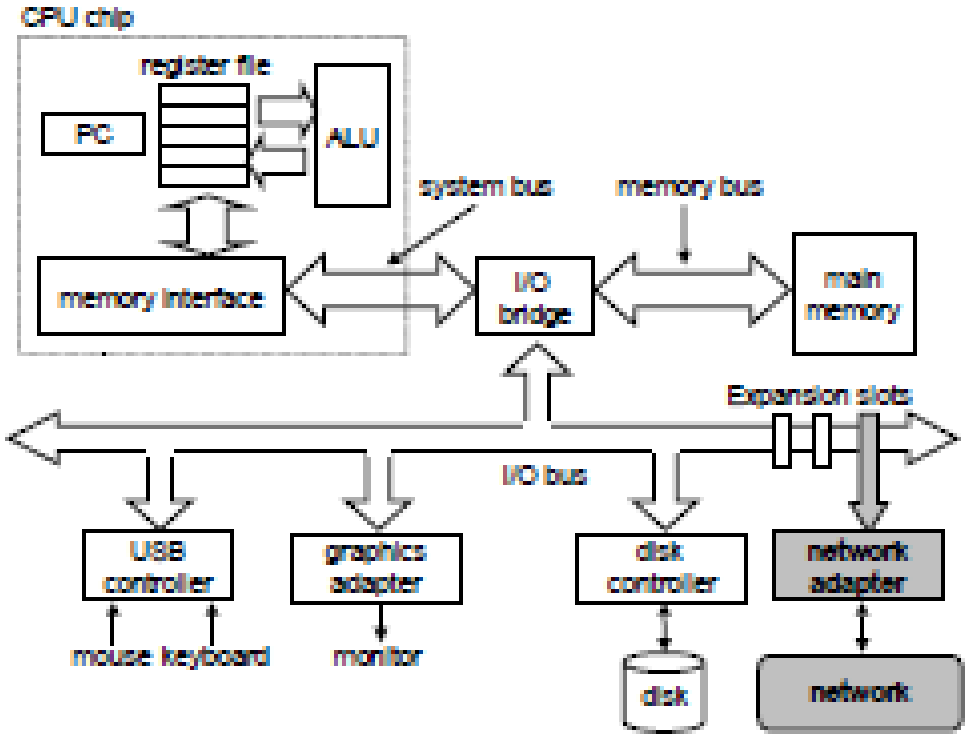

Networks

Remote execution (by telnet)

- 5 steps

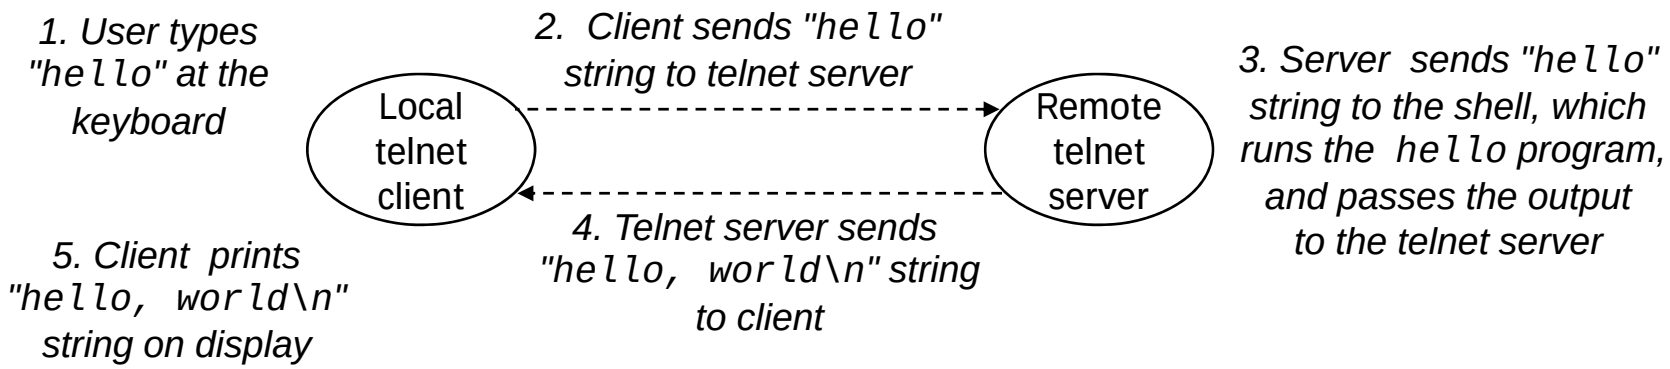

Concurrency and Parallelism

Concurrency

- Multiple processes execute at the same time, using processors alternately ☆
  - Allows multiple users to interact with the system at the same time
  - Allows a single user to engage in multiple tasks concurrently

Parallelism

- Supports multiple processors for the processes
- Multiprocessor systems
  - Multicore processors
  - Hyperthreading

Concurrency and Parallelism

■ Thread-level parallelism: Multicore processors

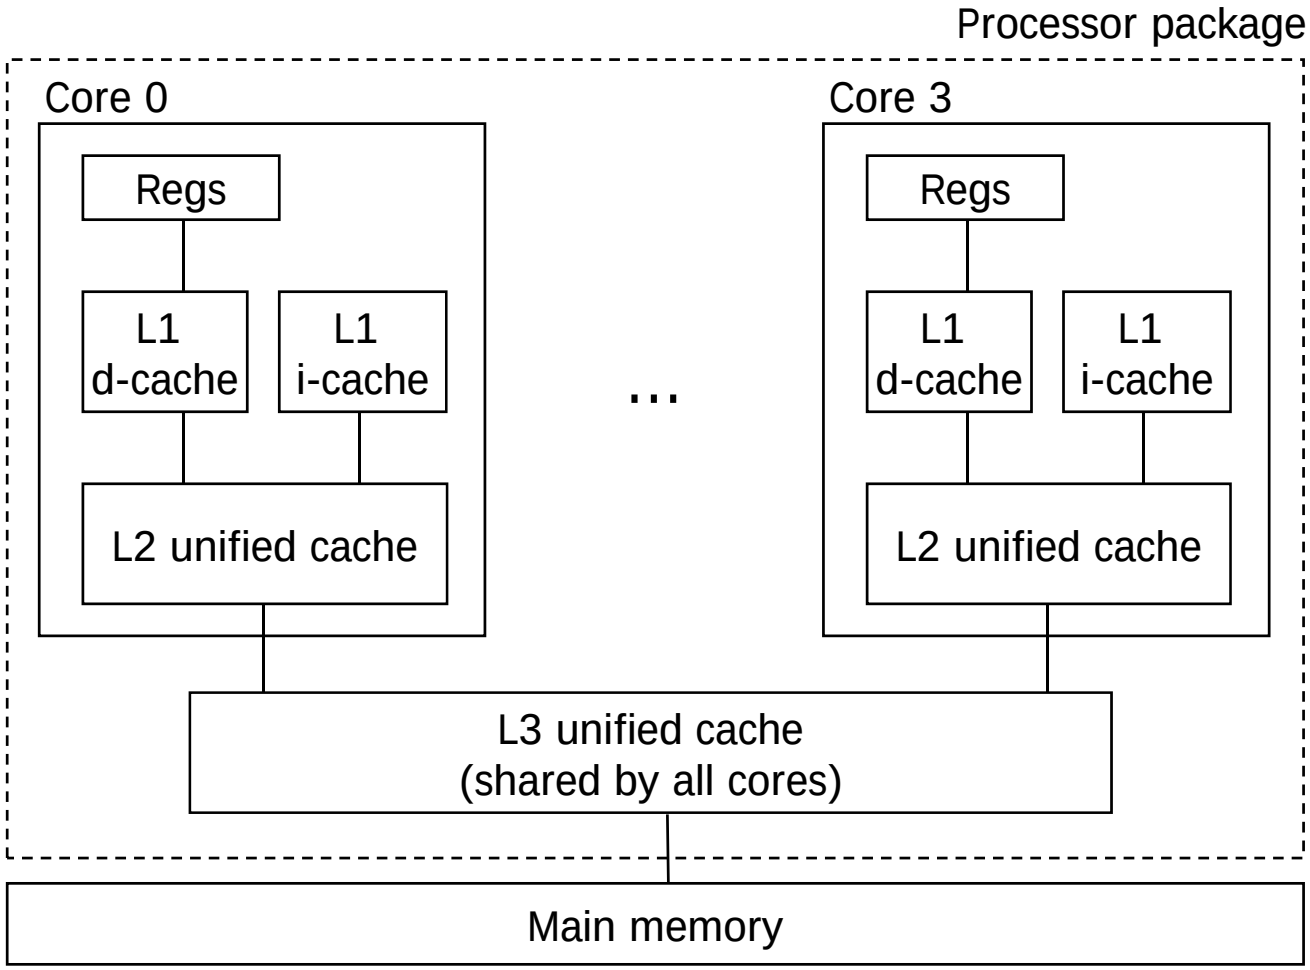

Concurrency and Parallelism

Thread-level parallelism: Hyperthreading

- A technique that allows a single CPU to execute multiple flows of control
  - Multiple copies of some of the CPU hardware (PCs and register files)
  - Only single copy of the other parts of the hardware (floating-point arithmetic units)
- Very short thread-switching time

- Intel Core i7
  - Supports 4 cores and 2 threads for each core
  - Can actually execute 8 threads in parallel

Concurrency and Parallelism

Instruction-level parallelism

- Pipelining
- Superscalar processors

SIMD parallelism

- Single-instruction multiple-data architecture
  - Allows a single instruction to cause multiple operations to be performed in parallel
  - In recent Intel/AMD processors
    - ✓ Add 8 single-precision floating-point numbers in parallel

Summary

backward-compatible