

[Chap.3-6] Machine-level Representation of Programs

Young Ik Eom (yieom@skku.edu, 031-290-7120)

Distributing Computing Laboratory

Sungkyunkwan University

<http://dclab.skku.ac.kr>



Contents

- ...
- Procedures
- **Compound data structures**
- Pointers
- GDB debugger
- Buffer overflow
- Floating-point codes

Compound Data Structures


■ Compound data structures

- Arrays
- Structures
- Unions
- Data alignment


Arrays

- Review: arrays and pointers in C
 - Row-major order memory allocation


```
int z1[3];
```




Programmer view



```
int z2[2][3];
```



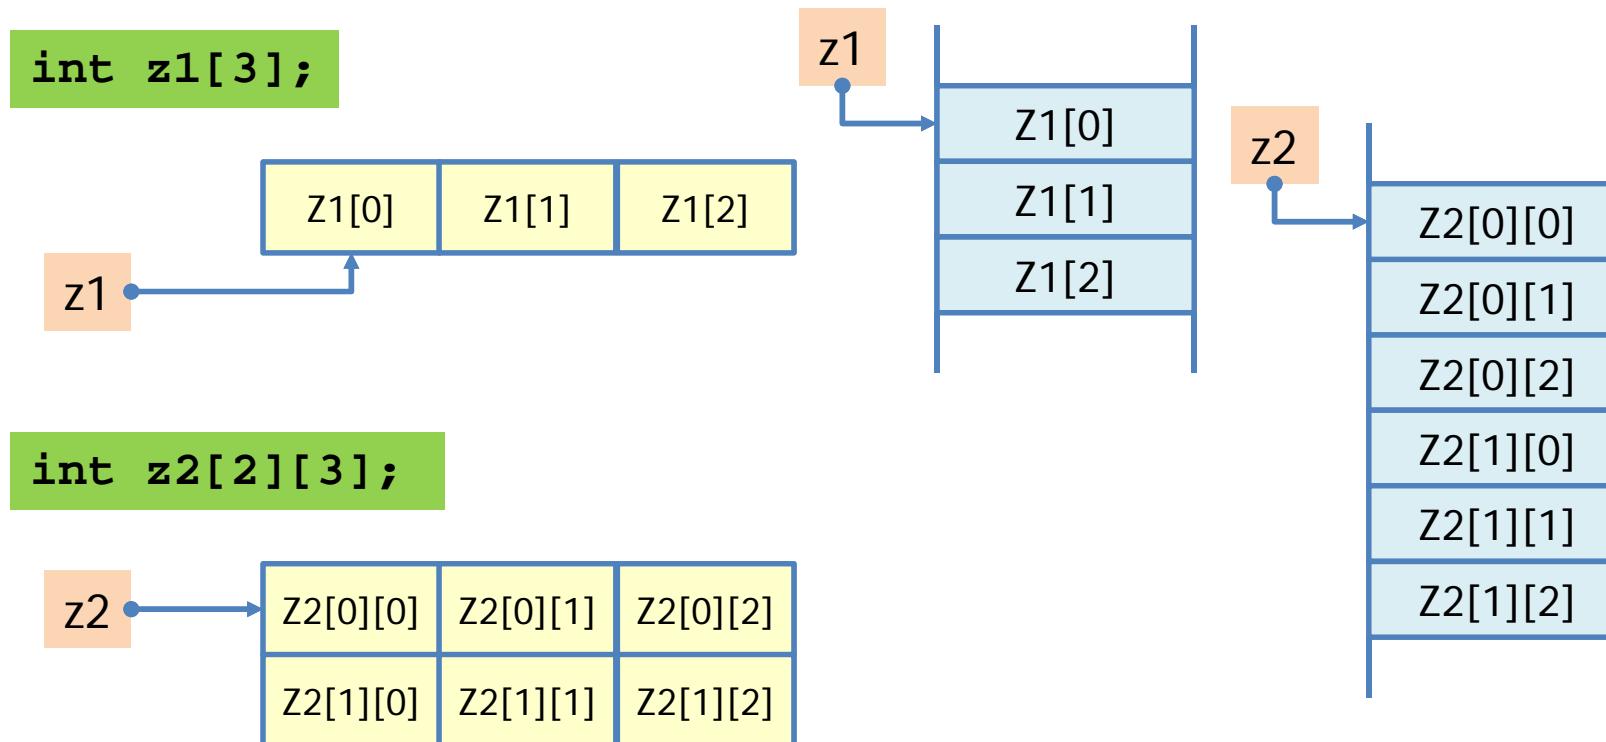
Programmer view



Arrays

■ Review: arrays and pointers in C

- Array name points to the 1st element of the array





Arrays

■ Review: arrays and pointers in C

▪ Pointer arithmetic

```
int z1[3];
```

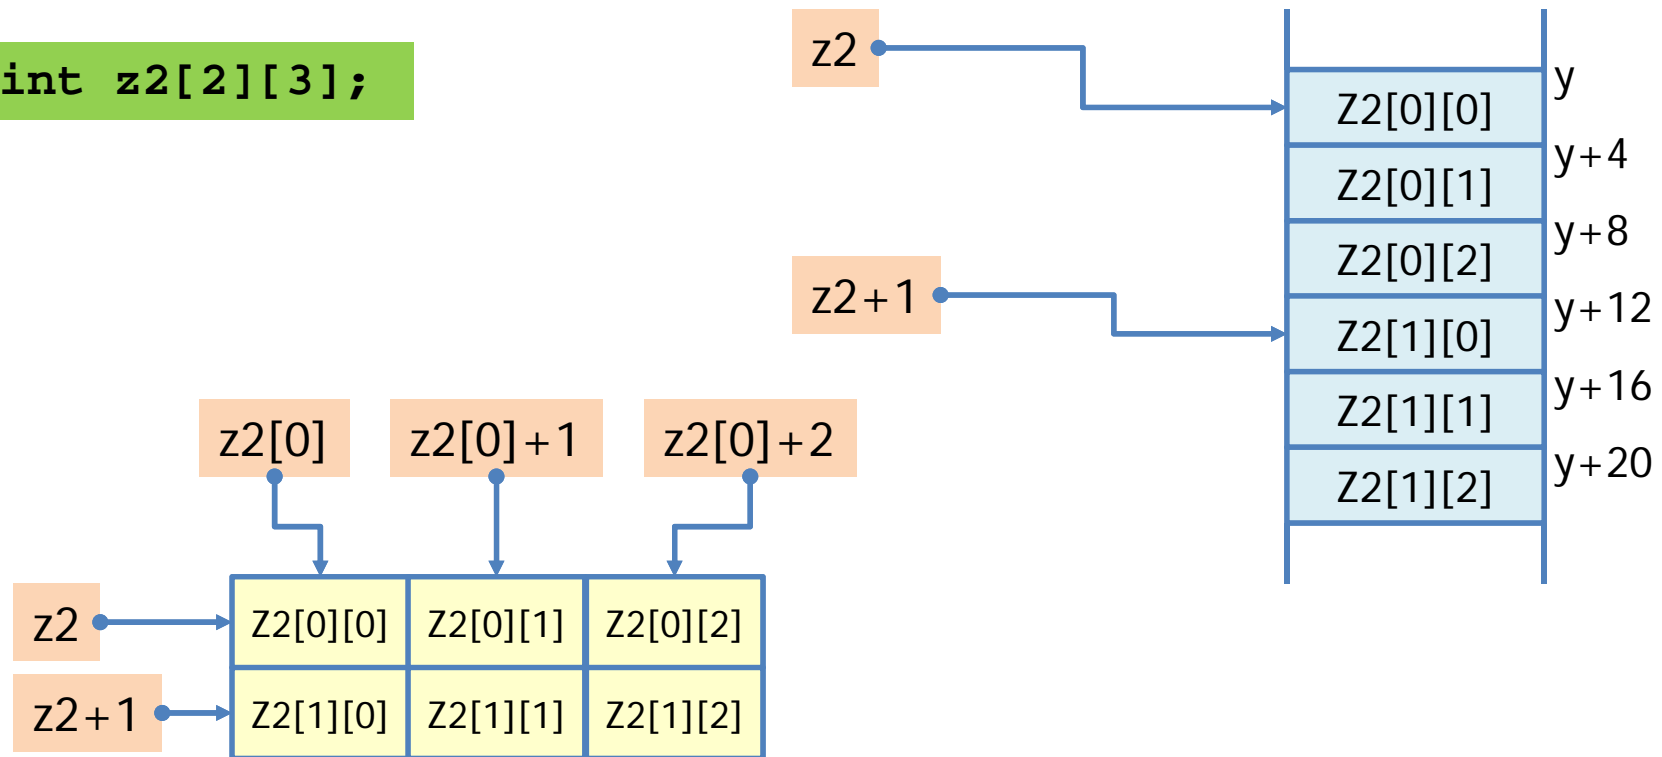


Arrays

■ Review: arrays and pointers in C

■ Pointer arithmetic

```
int z2[2][3];
```




Arrays

*char (*p) [3]*


■ Basic principles `T A[N];`

- Array of **N** elements of data type **T**
- Allocates a contiguous region of **N·sizeof(T)** bytes
- Array element **i** will be stored at address ($x_A + i \cdot \text{sizeof}(T)$)


`char s[12];`




`int v[5];`



`double x[4];`



`char *p[3];`




Arrays

■ Basic principles

- Memory referencing instructions of x86-64 are designed to simplify array access
- Example) `int z[10];`
 - The start address of the array `z` (x_z) in `%rdx`
 - Index `i` in `%rcx`
 - Then to copy `z[i]` into `%eax`

`movl (%rdx,%rcx,4),%eax`



Arrays

■ Pointer arithmetic

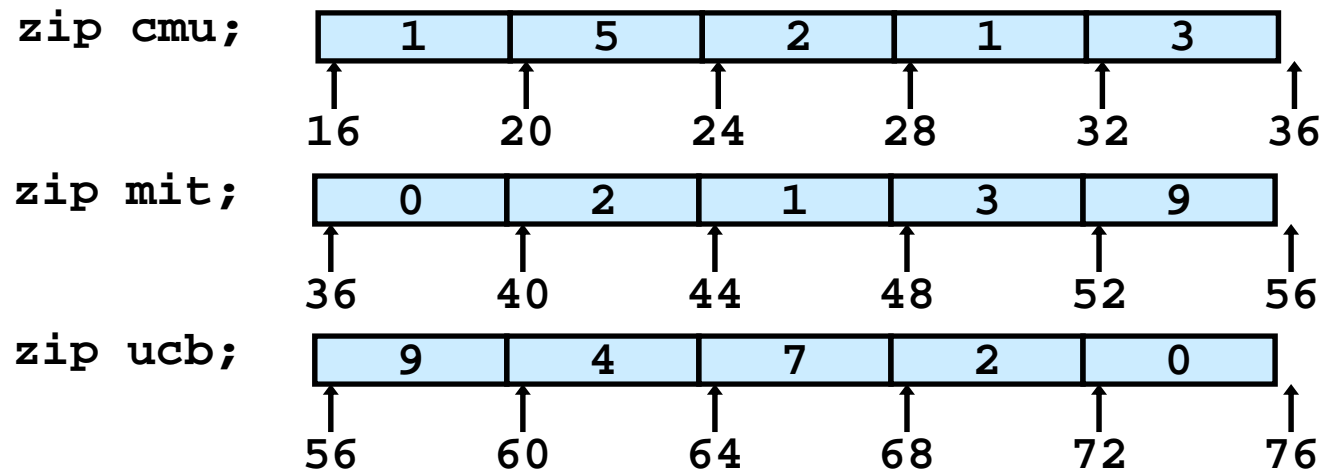
- The expression $\mathbf{p} + \mathbf{i}$ has value $\mathbf{x_p} + \mathbf{i \cdot sizeof(T)}$
 - \mathbf{p} is a pointer to data of type \mathbf{T}
 - The value of \mathbf{p} is $\mathbf{x_p}$
- Example) Array \mathbf{E} of integer
 - Starting address in $\mathbf{\%rdx}$ and index \mathbf{i} in $\mathbf{\%rcx}$
 - Destination $\mathbf{\%eax}$ or $\mathbf{\%rax}$

Expression	Type	Value	Assembly code
\mathbf{E}	$\mathbf{int^*}$	$\mathbf{x_E}$	$\mathbf{movq\ \%rdx,\ \%rax}$
$\mathbf{E[0]}$	\mathbf{int}	$\mathbf{M[x_E]}$	$\mathbf{movl\ (\%rdx),\ \%eax}$
$\mathbf{E[i]}$	\mathbf{int}	$\mathbf{M[x_E + 4i]}$	$\mathbf{movl\ (\%rdx,\ \%rcx, 4),\ \%eax}$
$\mathbf{\&E[2]}$	$\mathbf{int^*}$	$\mathbf{x_E + 8}$	$\mathbf{leaq\ 8(\%rdx),\ \%rax}$
$\mathbf{E + i - 1}$	$\mathbf{int^*}$	$\mathbf{x_E + 4i - 4}$	$\mathbf{leaq\ -4(\%rdx,\ \%rcx, 4),\ \%rax}$
$\mathbf{*(E + i - 3)}$	\mathbf{int}	$\mathbf{M[x_E + 4i - 12]}$	$\mathbf{movl\ -12(\%rdx,\ \%rcx, 4),\ \%eax}$
$\mathbf{\&E[i] - E}$	\mathbf{long}	\mathbf{i}	$\mathbf{movq\ \%rcx,\ \%rax}$

Arrays

■ Example)

```
typedef int zip[5];  
  
zip cmu = { 1, 5, 2, 1, 3 };  
zip mit = { 0, 2, 1, 3, 9 };  
zip ucb = { 9, 4, 7, 2, 0 };
```



■ Notes)


- Example arrays were allocated in successive 20 byte blocks
✓ Not guaranteed to happen in general

Arrays

■ Example) (Cont'd)

- No bounds checking in arrays!

Reference	Address	Value	Guaranteed?
<code>mit[3]</code>			
<code>mit[5]</code>			
<code>mit[-1]</code>			
<code>cmu[15]</code>			



Arrays

■ Example) Array & loop

```
typedef long zip[5];
...
long zip2int(zip z){
    int i;
    long zr = 0;
    for (i = 0; i < 5; i++)
        zr = 10 * zr + z[i];
    return zr;
}
```

Original source

$$zr = z[0] \cdot 10^4 + z[1] \cdot 10^3 + z[2] \cdot 10^2 + z[3] \cdot 10^1 + z[4]$$

- Array code to pointer code
- Eliminate loop variable i
- Use do-while loop structure

Transformed version

```
long zip2int(zip z){
    long zr = 0;
    long *zend = z + 4;
    do {
        zr = 10 * zr + *z;
        z++;
    } while(z <= zend);
    return zr;
}
```

Arrays

■ Example2) Array & loop

```
long zip2int(zip z){
    long zr = 0;
    long *zend = z + 4;
    do {
        zr = 10 * zr + *z;
        z++;
    } while(z <= zend);
    return zr;
}
```

[Registers]

%rdi	z
%rax	zr
%rbx	zend

z++
increments
by 8

$10 * zr + *z$
 $= *z + 2*(zr+4*zr)$

z in %rdi

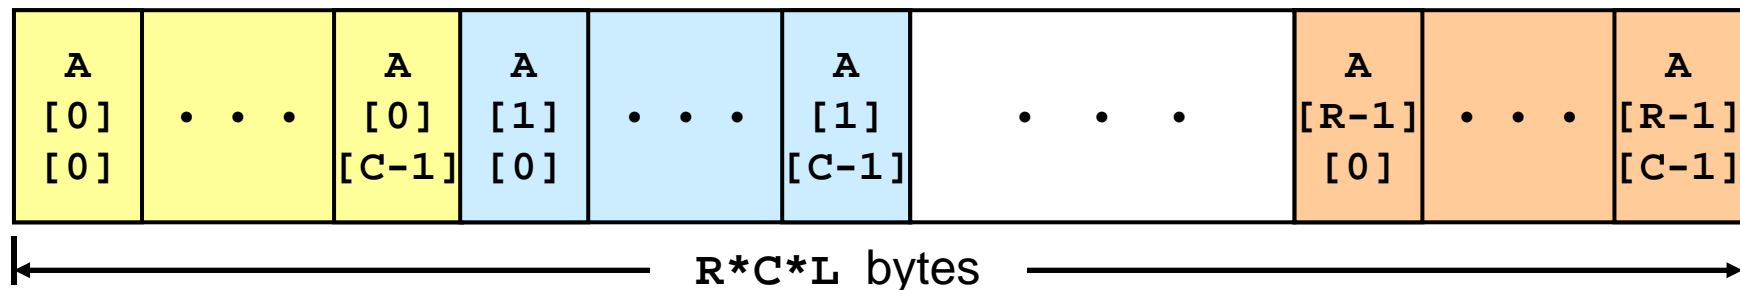
```
xorq %rax,%rax
leaq 32(%rdi),%rbx
.L59:
leaq (%rax,%rax,4),%rdx
movq (%rdi),%rax
addq $8,%rdi
leaq (%rax,%rdx,2),%rax
cmpq %rbx,%rdi
jle .L59
ret
```

```
zr = 0
zend = z + 4

Compute 5*zr
Access *z
Compute z++
zr = *z + 2*(5*zr)
Compare z : zend
if <=, goto loop
```

```
T A[R][C];
```

- $$\begin{bmatrix} A[0][0] & \bullet & \bullet & \bullet & A[0][C-1] \\ & \bullet & & & \bullet \\ & \bullet & & & \bullet \\ & \bullet & & & \bullet \\ A[R-1][0] & \bullet & \bullet & \bullet & A[R-1][C-1] \end{bmatrix}$$



Arrays


■ Nested arrays

```
T A[R][C];
```

- Variable A denotes an array of R elements, where each element is an array of C elements
- Example)

```
int pgh[4][5] =  
    {{1, 5, 2, 0, 6},  
     {1, 5, 2, 1, 3},  
     {1, 5, 2, 1, 7},  
     {1, 5, 2, 2, 1}};
```

- Variable **pgh** denotes an array of 4 elements (allocated contiguously)
- Each element is an array of 5 integers (allocated contiguously)

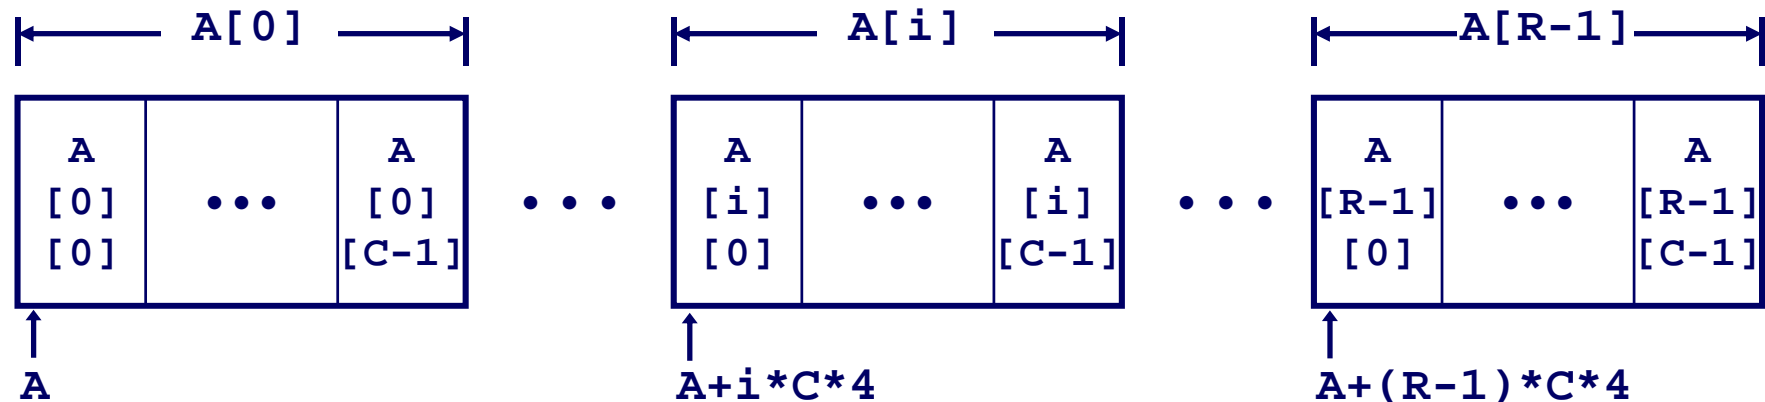


Arrays

■ Nested arrays: Row vectors `T A[R][C];`

- $A[i]$ is an array of C elements
- Starting address of $A[i]$ at $A + i \cdot (C \cdot L)$

```
int A[R][C];
```



Arrays

■ Nested arrays: Row vectors

■ Example) `int pgh[4][5];`

- `pgh[i]` is an array of 5 integers
- Getting the starting address of `pgh[i]`
 - ✓ `pgh + 20*i`

```
int *get_pgh_zip(int i)
{
    return pgh[i];
}
```

- Compute as `pgh + 4 * (i + 4 * i)`

`i in %rdi`

```
leaq (%rdi,%rdi,4),%rax
leaq pgh(,%rax,4),%rax
ret
```

Compute `5*i`


Compute `pgh+(20*i)`

Arrays

■ Nested arrays: Array elements `T A[R][C];`

- `A[i][j]` is an element of type `T`
 - Address at $A + i \cdot (C \cdot L) + j \cdot L = A + (i \cdot C + j) \cdot L$

```
int A[R][C];
```



Arrays

■ Nested arrays: Array elements

■ Example) `int pgh[4][5];`

- `pgh[i][j]` is an integer
- Getting the address of `pgh[i][j]`
 - ✓ $\text{pgh} + 20 \cdot i + 4 \cdot j$

```
int get_pgh_digit(int i, int j)
{
    return pgh[i][j];
}
```

- Compute as $\text{pgh} + 4 * j + 4 * (i + 4 * i)$

`i in rdi, j in rsi`

`leaq 0(,%rsi,4),%rdx`

Compute $4*j$

`leaq (%rdi,%rdi,4),%rax`

Compute $5*i$

`movl pgh(%rdx,%rax,4),%eax`


Access $*(\text{pgh} + 4*j + 20*i)$

`ret`

Arrays

■ Nested arrays: Strange referencing examples

```
int pgh[4][5];
```



Reference	Address	Value	Guaranteed?
<code>pgh[3][3]</code>			
<code>pgh[2][5]</code>			
<code>pgh[2][-1]</code>			
<code>pgh[4][-1]</code>			
<code>pgh[0][19]</code>			
<code>pgh[0][-1]</code>			

Arrays

■ Fixed-size arrays

- Optimizations by gcc
 - Translation of array expressions into pointer expressions
- Example)
 - Compute element i,k of the product of matrices A and B
 - $\sum_{j=0}^{n-1} a_{i,j} \cdot b_{j,k}$

```
#define N 16  
typedef int f_mat[N][N];
```

Arrays

■ Fixed-size arrays

■ Example)

[Original C code]

```
/* Compute element i,k of fixed matrix product */

int fm_pr_el(f_mat A, f_mat B, long i, long k)
{
    long j;
    int res = 0;
    for (j = 0; j < N; j++)
        res += A[i][j]*B[j][k];
    return res;
}
```

Arrays

■ Fixed-size arrays

■ Example)

[Optimized C code]

```
/* Compute element i,k of fixed matrix product */

int fm_pr_el_opt(f_mat A, f_mat B, long i, long k)
{
    int *Aptr = &A[i][0];
    int *Bptr = &B[0][k];
    int *Bend = &B[N][k];
    int res = 0;
    do {
        res += *Aptr * *Bptr;
        Aptr++;
        Bptr += N;
    } while (Bptr != Bend);
    return res;
}
```


Arrays

■ Fixed-size arrays

■ Example)

[Assembly code]

```
int fm_pr_el(f_mat A, f_mat B, long i, long k)
A in %rdi, B in %rsi, i in %rdx, k in %rcx
fm_pr_el:
    salq $6,%rdx           Compute 64*i
    addq %rdx,%rdi         Compute Aptr = A+64*i = &A[i][0]
    leaq (%rsi,%rcx,4),%rcx Compute Bptr = B+4k = &B[0][k]
    leaq 1024(%rcx),%rsi    Compute Bend = B+4k+1024 = &B[N][k]
    movl $0,%eax           Set res = 0
.L7:
    movl (%rdi),%edx        Read *Aptr
    imull (%rcx),%edx       Multiply by *Bptr
    addl %edx,%eax          Add to res
    addq $4,%rdi            Increment Aptr++
    addq $64,%rcx           Increment Bptr += N
    cmpq %rsi,%rcx         Compare Bptr:Bend
    jne .L7                If !=, goto loop
    rep; ret               Return
```

Arrays

■ Variable-size arrays

- Historically, C only supported multi-dimensional arrays where the sizes could be determined at compile time
 - With the possible exception of the first dimension
- ISO C99 introduced the capability of having array dimension expressions that are computed as the array is being allocated
- Declaration of variable-size array
 - Either as a local variable or function argument

```
int A[expr1][expr2];
```

- The dimensions of the array are determined by evaluating the expressions **expr1** and **expr2** at the time the declaration is encountered

Arrays

■ Variable-size arrays

■ Example)

```
int vm_el(long n, int A[n][n], long i, long j)
{
    return A[i][j];
}
```

```
int vm_el(long n, int A[n][n], long i, long j)
n in %rdi, A in %rsi, i in %rdx, j in %rcx

vm_el:
    imulq %rdx,%rdi           Compute n*i
    leaq (%rsi,%rdi,4),%rax    Compute A + 4*n*i
    movl (%rax,%rcx,4),%eax    Read from M[A + 4*n*i + 4*j]
    ret                       Return
```

Arrays

■ Summary

- Arrays in C
 - Contiguous allocation of memory
 - Same type of all elements
 - Array name
 - ✓ Pointer (constant) to the first element
 - No bounds checking
- Compiler optimizations
 - Compiler often turns array code into pointer code
 - Uses addressing modes to scale array indices
 - Lots of tricks to improve array indexing in loops

Structures

■ Structure

- Aggregates multiple objects of different types into a single unit
 - Contiguous allocation of memory
 - Members may be of different types
 - Refer to members by name (in C)
 - Refer to members by offset (in machine code)
- Uses keyword **struct** in C
 - No structure in machine code

Structures


■ Structure

- Member access
 - By name in C and by offset in assembly/machine code

```
struct rec {  
    int i;  
    int j;  
    int a[2];  
    int *p;  
};
```

```
void set(struct rec *r)  
{  
    r->j = r->i;  
}
```

Memory Layout



Assembly code

```
r in %rdi  
  
set:  
    movl (%rdi),%eax    Get r->i  
    movl %eax,4(%rdi)   Store in r->j  
    ret
```

Structures


■ Structure

- Getting pointer to a member
 - Offset of each member determined at compile time

```
struct rec {  
    int i;  
    int j;  
    int a[2];  
    int *p;  
};
```

```
int *get(struct rec *r,  
        int i)  
{  
    return &(r->a[i]);  
}
```

Memory Layout



Assembly code

```
r in %rdi, i in %rsi
```

```
get:
```

```
leaq 8(%rdi,%rsi,4),%rax &(r->a[i])  
ret
```

Structures


■ Structure

- Getting pointer to a member
 - Offset of each member determined at compile time

```
struct rec {  
    int i;  
    int j;  
    int a[2];  
    int *p;  
};
```

```
void put(struct rec *r)  
{  
    r->p = &r->a[r->i  
          + r->j];  
}
```

Memory Layout



Assembly code

```
r in %rdi  
put:  
    movl 4(%rdi),%eax  
    addl (%rdi),%eax  
    cltq  
    leaq 8(%rdi,%rax,4),%rax  
    movq %rax,16(%rdi)
```


Data Alignment



■ Alignment restrictions

- Restrictions on the allowable addresses for the primitive data types
 - Primitive data type that requires **K** bytes
 - ✓ Its address must be multiple of **K**
 - Required on many machines
 - Recommended on x86-64
- Simplify the design of hardware
(Interface between the processor and the memory)
- Improve the memory system performance

Data Alignment

■ Alignment restrictions

- 1 byte (char)
 - No restrictions on address
- 2 bytes (short)
 - Must have address that is a multiple of 2
- 4 bytes (int, float)
 - Must have address that is a multiple of 4
- 8 bytes (long, double, char *, ...)
 - Must have address that is a multiple of 8

Data Alignment

■ Compiler directives for alignment

- The compiler places directives in the assembly code indicating the desired alignment for data
- Example)

.align 8

- Align address to multiple of 8

```
.section .rodata
.align 8
.L4:
    .quad .L3          Case 100: l_A
    .quad .L8          Case 101: l_def
    .quad .L5          Case 102: l_B
    .quad .L6          Case 103: l_C
    .quad .L7          Case 104: l_D
    .quad .L8          Case 105: l_def
    .quad .L7          Case 106: l_D
```

Data Alignment

■ Alignments for structures

- Member alignments for structures
 - Member offset must satisfy the element's alignment requirement
- Overall structure alignment
 - Each structure has alignment requirement **K**
 - ✓ **K** is the largest alignment requirement of the structure elements
 - ✓ Initial address & structure length must be multiples of **K**

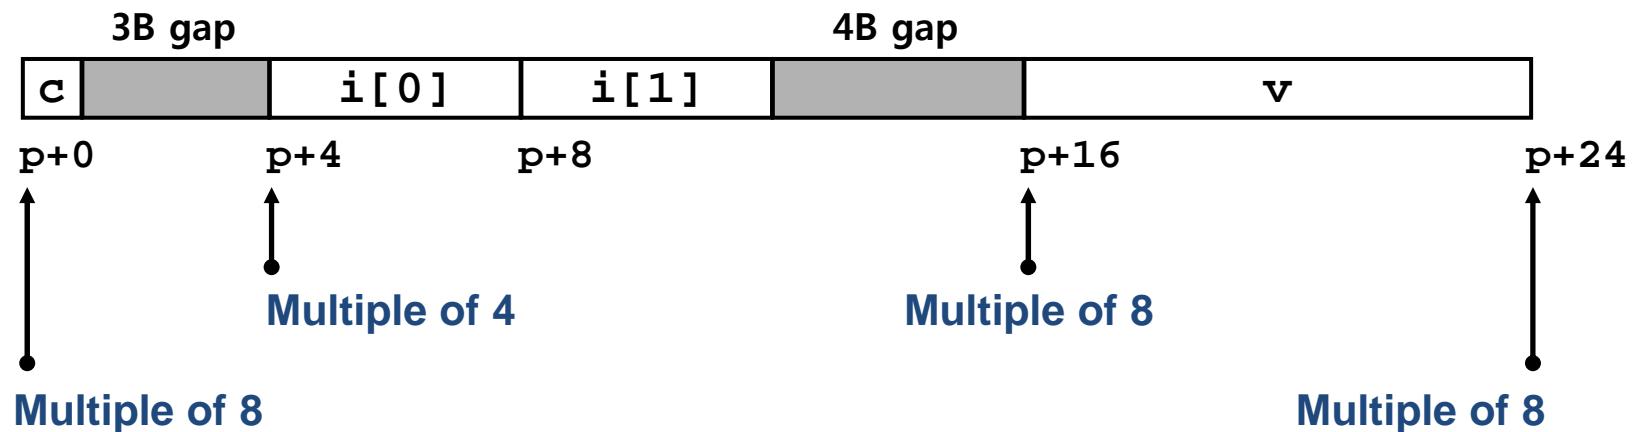
Data Alignment

■ Alignments for structures: Member alignment

■ Example)

- Gap insertion for member alignment
- Structure alignment $K = 8$
(due to `double` elements)

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```




Data Alignment

■ Alignments for structures: Overall alignment


```
struct S2 {  
    double x;  
    int i[2];  
    char c;  
} *p;
```

p must be multiple of 8



```
struct S3 {  
    float x[2];  
    int i[2];  
    char c;  
} *p;
```

p must be multiple of 4



Data Alignment

■ Alignments for structures: Ordering elements


```
struct s4 {  
    char c1;  
    double v;  
    char c2;  
    int i;  
} *p;
```

10 bytes wasted space



```
struct s5 {  
    double v;  
    char c1;  
    char c2;  
    int i;  
} *p;
```

2 bytes wasted space



Unions


■ Principles

- Allows an object to be referenced using several different types
 - Overlays union elements
 - Allocates space according to largest elements
 - Usage
 - ✓ Reducing space allocations
 - ✓ Accessing the bit patterns of different data types
- Uses keyword **union** in C


Unions

■ Union example

```
union U1 {  
    char c;  
    int i[2];  
    double d;  
} *p;
```



```
struct S1 {  
    char c;  
    int i[2];  
    double d;  
} *q;
```



Type	Offset (c)	Offset (i)	Offset (d)	Size
U1	0	0	0	8
S1	0	4	16	24


Unions

■ Union

■ Byte ordering

```
double b2d (unsigned w0, unsigned w1)
{
    union U1 {
        double d;
        unsigned u[2];
    } temp;

    temp.u[0] = w0;
    temp.u[1] = w1;
    return temp.d;
}
```



- On little-endian machine, (x86-64)
 - **w0** becomes the low-order 4 bytes of **d**
 - **w1** becomes the high-order 4 bytes of **d**
- On big-endian machine,
 - **w0** becomes the high-order 4 bytes of **d**
 - **w1** becomes the low-order 4 bytes of **d**

Structures and Unions



■ Summary

■ Structure

- Allocates elements in the order declared
- Padding in the middle and at end to satisfy alignment

■ Union

- Overlays elements
- Way to circumvent type system

Summary

- Array
- Structure
- Union