

[Chap.6-4] The Memory Hierarchy

Young Ik Eom (yieom@skku.edu, 031-290-7120)
Distributing Computing Laboratory
Sungkyunkwan University
<http://dclab.skku.ac.kr>



Contents

- Storage technologies
- Locality
- Memory hierarchy
- Cache memories
- Writing cache-friendly code
- Impact of caches on program performance


Example: Rearranging Loops

■ Rearranging loops to increase spatial locality (Matrix multiplication)

■ Assumptions

- $n \times n$ array of double, $\text{sizeof}(\text{double}) = 8$
- Single cache with 32B block size
- The array size n is so large
(Single matrix row does not fit in the cache)
- Local variables in registers


■ Access pattern



Example: Rearranging Loops

■ Ordinary code for matrix multiplication

```
/* matrix multiplication */  
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        c[i][j] = 0.0;  
        for (k=0; k<n; k++)  
            c[i][j] += a[i][k] * b[k][j];  
    }  
}
```




```
c[i][j] = c[i][j] + a[i][k] * b[k][j];
```

Example: Rearranging Loops

■ Rearranging loops to increase spatial locality (Matrix multiplication: version-ijk)

```
/* ijk */  
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```




```
c[i][j] = c[i][j] + a[i][k] * b[k][j];
```

- Miss rates per inner loop iteration
 - a (0.25), b (1.00), c (0.00)

Example: Rearranging Loops

■ Rearranging loops to increase spatial locality (Matrix multiplication: version-jik)

```
/* jik */
for (j=0; j<n; j++) {
    for (i=0; i<n; i++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}
```



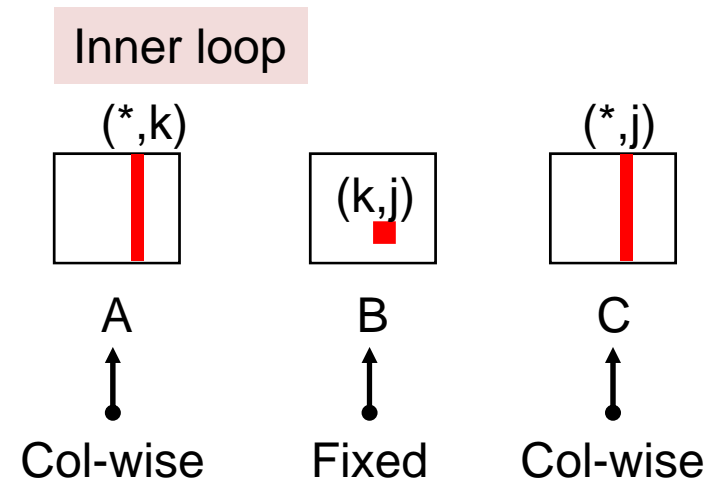
```
c[i][j] = c[i][j] + a[i][k] * b[k][j];
```

- Miss rates per inner loop iteration
 - a (0.25), b (1.00), c (0.00)

Example: Rearranging Loops

■ Rearranging loops to increase spatial locality (Matrix multiplication: version-jki)

```
/* jki */  
for (j=0; j<n; j++) {  
    for (k=0; k<n; k++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```



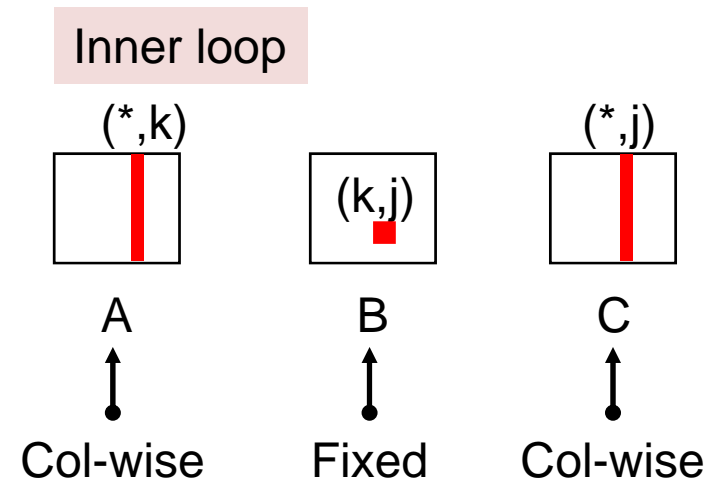
```
c[i][j] = c[i][j] + a[i][k] * b[k][j];
```

- Miss rates per inner loop iteration
 - a (1.00), b (0.00), c (1.00)

Example: Rearranging Loops

■ Rearranging loops to increase spatial locality (Matrix multiplication: version-kji)

```
/* kji */  
for (k=0; k<n; k++) {  
    for (j=0; j<n; j++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```



```
c[i][j] = c[i][j] + a[i][k] * b[k][j];
```

- Miss rates per inner loop iteration
 - a (1.00), b (0.00), c (1.00)


Example: Rearranging Loops

■ Rearranging loops to increase spatial locality (Matrix multiplication: version-kij)

```
/* kij */  
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

```
c[i][j] = c[i][j] + a[i][k] * b[k][j];
```

- Miss rates per inner loop iteration
 - a (0.00), b (0.25), c (0.25)




Example: Rearranging Loops

■ Rearranging loops to increase spatial locality (Matrix multiplication: version-ikj)

```
/* ikj */  
for (i=0; i<n; i++) {  
    for (k=0; k<n; k++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

```
c[i][j] = c[i][j] + a[i][k] * b[k][j];
```

- Miss rates per inner loop iteration
 - a (0.00), b (0.25), c (0.25)



Example: Rearranging Loops

■ Rearranging loops to increase spatial locality (Matrix multiplication: summary)

```
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```

```
for (j=0; j<n; j++) {  
    for (k=0; k<n; k++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```

```
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```


- ijk (& jik):
 - 2 loads, 0 stores
 - miss-ratio/iter = 1.25

- jki (& kji):
 - 2 loads, 1 store
 - miss-ratio/iter = 2.0

- kij (& ikj):
 - 2 loads, 1 store
 - miss-ratio/iter = 0.5

Example: Rearranging Loops

- Rearranging loops to increase spatial locality
(Matrix multiplication: summary)



Example: Rearranging Loops

- **Rearranging loops to increase spatial locality (Matrix multiplication: some points)**
 - For large values of n , the fastest version runs almost 40 times faster than the slowest version, even though they perform the same # of FP operations
 - Miss ratio is a better predictor of performance than the total # of memory references (in this case)
 - For large values of n , the performance of the fastest pair of versions is constant
 - Etc.

Exploiting Locality in Pgms

- Programs with good locality
 - Access most of their data from fast cache memories
 - Programs with poor locality
 - Access most of their data from slow DRAM memories
- Programmers must exploit this to write more efficient programs

Tips

- Focus your attention on the inner loops
- Try to maximize the spatial locality in your programs
 - Reading data objects sequentially, with stride-1
- Try to maximize the temporal locality in your programs
 - Using data objects repeatedly once they are read from memory

Summary

